

# 01\_LG(TakeHome)

August 23, 2021

## 0.0.1 Take-home exercise (40 points)

Find an interesting dataset for linear regression on Kaggle. Implement the normal equations and gradient descent then evaluate your model's performance.

Write a brief report on your experiments and results in the form of a Jupyter notebook.

Explain the dataset which you get and which rows which you use. How many data in your dataset?

Write down your all code at below. Show the results, goodness of fit and plot cost graph

```
[1]: NAME = "Kanawut Kaewnoparat"
      ID = "st122109"
```

```
[2]: import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      import seaborn as sns
```

```
[3]: data = pd.read_csv("insurance.csv")
```

```
[4]: data.head()
```

```
[4]:   age    sex    bmi  children  smoker    region    charges
0    19  female  27.900         0     yes  southwest  16884.92400
1    18   male  33.770         1     no   southeast   1725.55230
2    28   male  33.000         3     no   southeast   4449.46200
3    33   male  22.705         0     no  northwest  21984.47061
4    32   male  28.880         0     no  northwest   3866.85520
```

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column    Non-Null Count  Dtype
---  -
0   age       1338 non-null   int64
1   sex       1338 non-null   object
2   bmi       1338 non-null   float64
```

```
3  children  1338 non-null  int64
4  smoker    1338 non-null  object
5  region    1338 non-null  object
6  charges   1338 non-null  float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
[6]: data.dtypes
```

```
[6]: age          int64
sex           object
bmi          float64
children     int64
smoker       object
region       object
charges      float64
dtype: object
```

```
[7]: data.isnull().sum()
```

```
[7]: age          0
sex          0
bmi          0
children     0
smoker       0
region       0
charges      0
dtype: int64
```

```
[8]: data.select_dtypes(include = [int, float]).head()
```

```
[8]:   age    bmi  children    charges
0   19  27.900         0  16884.92400
1   18  33.770         1   1725.55230
2   28  33.000         3   4449.46200
3   33  22.705         0  21984.47061
4   32  28.880         0   3866.85520
```

```
[9]: data.select_dtypes(exclude= [int, float]).columns
```

```
[9]: Index(['sex', 'smoker', 'region'], dtype='object')
```

```
[10]: data['sex'].unique()
```

```
[10]: array(['female', 'male'], dtype=object)
```

```
[11]: categorical_col = data.select_dtypes(exclude= [int, float]).columns
      for i in categorical_col:
          print(f"Unique valeus in {i}: {data[i].unique()}")
      print(categorical_col)
```

```
Unique valeus in sex: ['female' 'male']
Unique valeus in smoker: ['yes' 'no']
Unique valeus in region: ['southwest' 'southeast' 'northwest' 'northeast']
Index(['sex', 'smoker', 'region'], dtype='object')
```

```
[12]: numerical_col = data.select_dtypes(include= [int, float]).columns
      print(numerical_col)
```


```
Index(['age', 'bmi', 'children', 'charges'], dtype='object')
```

```
[13]: for index, value in enumerate(numerical_col):
      print(index, value)
```

```
0 age
1 bmi
2 children
3 charges
```

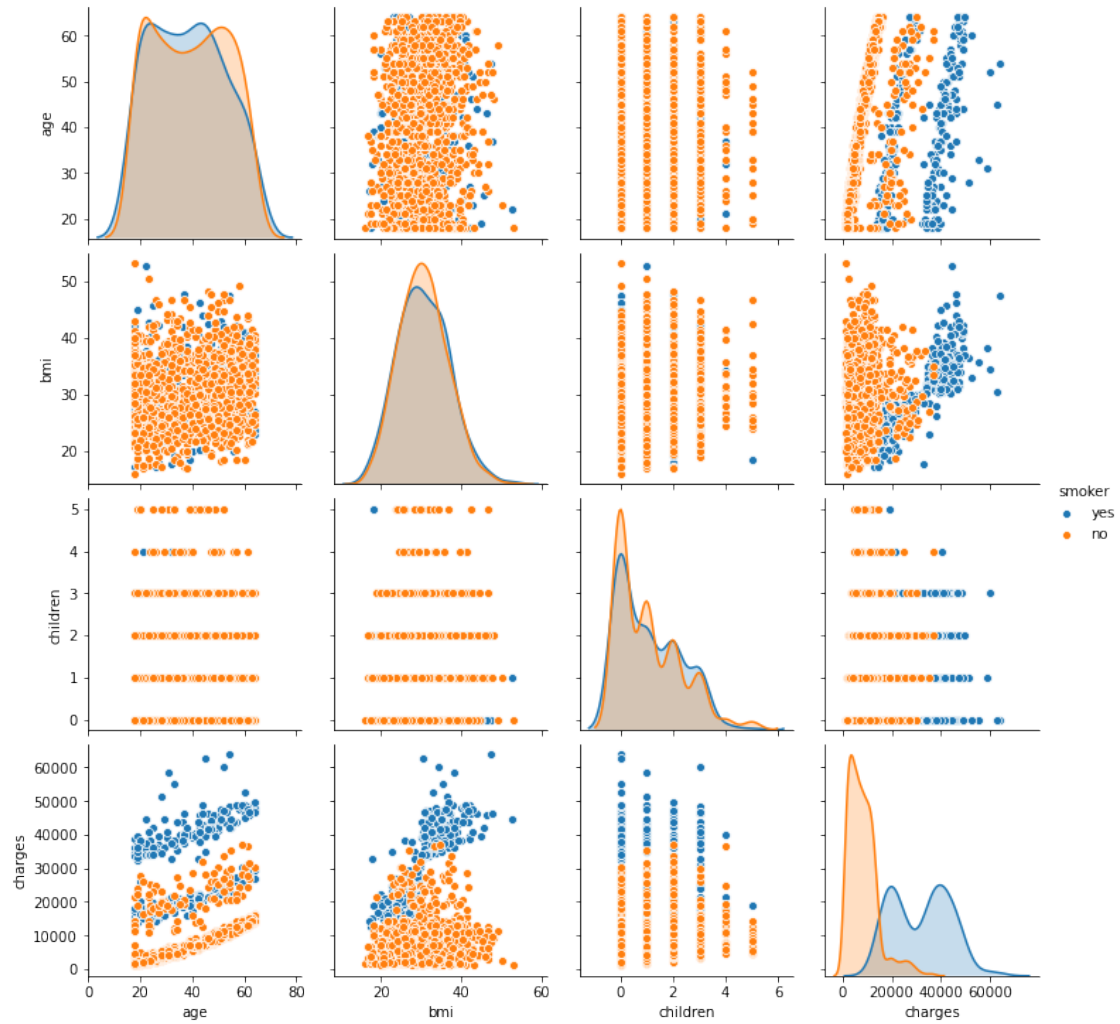
```
[14]: fig, axes = plt.subplots(2,2, figsize = (15,10))
      sns.distplot(a = data['age'], kde = False, ax = axes[0,0] )
      sns.distplot(a = data['bmi'], kde = False, ax = axes[0,1] )
      sns.distplot(a = data['children'], kde = False, ax = axes[1,0] )
      sns.distplot(a = data['charges'], kde = False, ax = axes[1,1] )
```

```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7f186f2c9b20>
```



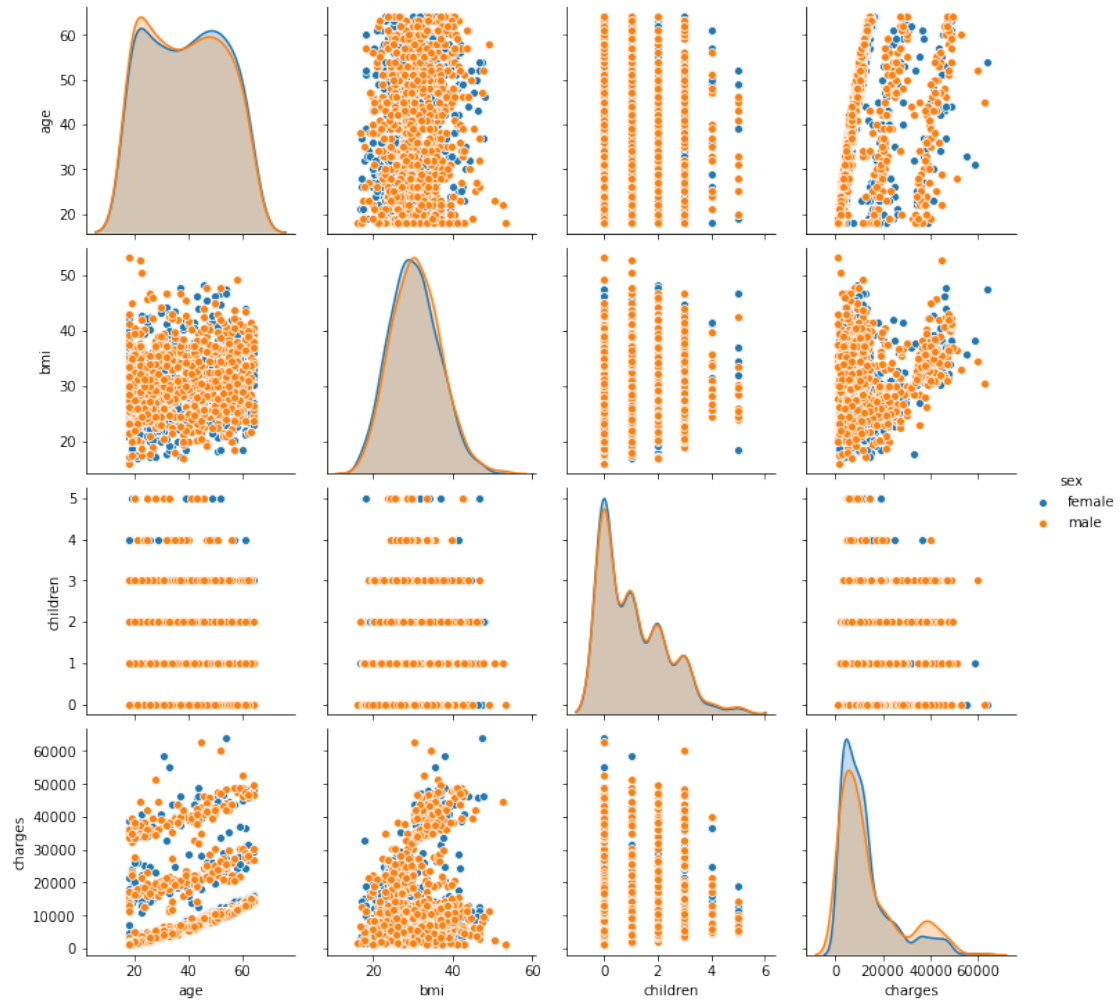
```
[15]: sns.pairplot(data, hue = 'smoker')
```

```
[15]: <seaborn.axisgrid.PairGrid at 0x7f186ee50f10>
```



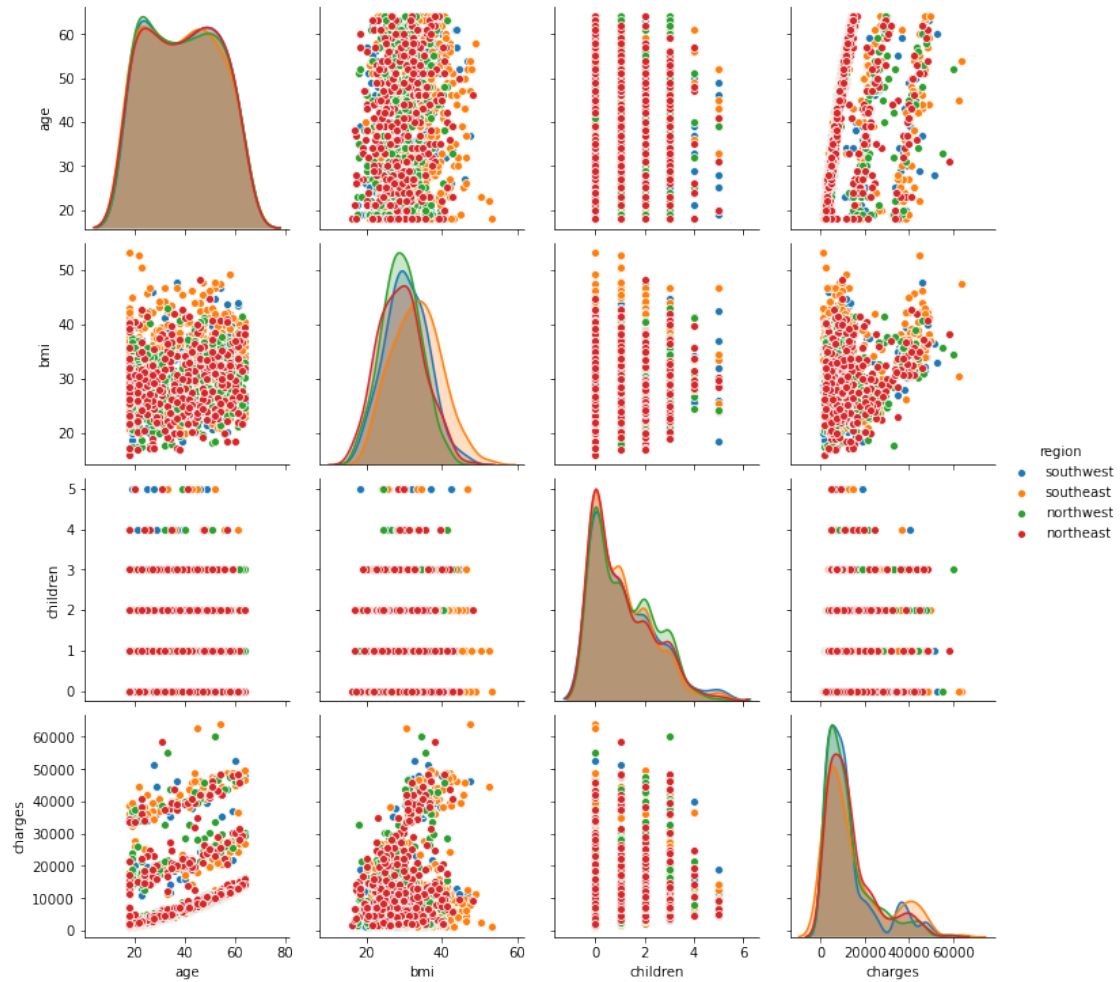
```
[16]: sns.pairplot(data, hue = 'sex')
```

```
[16]: <seaborn.axisgrid.PairGrid at 0x7f186e920a30>
```




```
[17]: sns.pairplot(data, hue = 'region')
```

```
[17]: <seaborn.axisgrid.PairGrid at 0x7f186cb13d90>
```



```
[18]: grid = sns.FacetGrid(data= data, col = 'smoker', hue = 'sex')
      grid.map(sns.scatterplot, 'age', 'charges')
```

```
[18]: <seaborn.axisgrid.FacetGrid at 0x7f186e978190>
```



```
[19]: data[numerical_col].corr(method = 'kendall')
```


```
[19]:
```

	age	bmi	children	charges
age	1.000000	0.073273	0.043253	0.475302
bmi	0.073273	1.000000	0.011562	0.082524
children	0.043253	0.011562	1.000000	0.103107
charges	0.475302	0.082524	0.103107	1.000000

```
[20]: sns.heatmap(data[numerical_col].corr(method = 'kendall'), annot = True)
```


```
[20]: <matplotlib.axes._subplots.AxesSubplot at 0x7f186c463c40>
```






```
[21]: sns.lmplot(x = 'age', y='charges', data =data, ci = 0.95, hue = 'smoker')
```

```
[21]: <seaborn.axisgrid.FacetGrid at 0x7f186c376190>
```




```
[22]: sns.boxplot(x = 'smoker', y='charges', data =data)
```

```
[22]: <matplotlib.axes._subplots.AxesSubplot at 0x7f186c2f37c0>
```




```
[23]: sns.lmplot(x = 'bmi', y = 'charges', hue = 'smoker', col = 'smoker', data = data)
```

```
[23]: <seaborn.axisgrid.FacetGrid at 0x7f186c255520>
```



```
[24]: sns.boxplot(x = 'region', y = 'charges', data = data, hue = 'smoker')
```

```
[24]: <matplotlib.axes._subplots.AxesSubplot at 0x7f186c1e62e0>
```



```
[25]: data[data['smoker'] == 'no'].describe()
```

```
[25]:
```

	age	bmi	children	charges
count	1064.000000	1064.000000	1064.000000	1064.000000
mean	39.385338	30.651795	1.090226	8434.268298
std	14.083410	6.043111	1.218136	5993.781819
min	18.000000	15.960000	0.000000	1121.873900
25%	26.750000	26.315000	0.000000	3986.438700
50%	40.000000	30.352500	1.000000	7345.405300
75%	52.000000	34.430000	2.000000	11362.887050
max	64.000000	53.130000	5.000000	36910.608030

```
[26]: data['region'].value_counts(normalize = True)
```

```
[26]:
```

southeast	0.272048
northwest	0.242900
southwest	0.242900
northeast	0.242152

Name: region, dtype: float64

```
[27]: data['sex'].value_counts(normalize = True)
```

```
[27]: male      0.505232
      female    0.494768
      Name: sex, dtype: float64
```

```
[28]: data['smoker'].value_counts(normalize = True)
```

```
[28]: no      0.795217
      yes     0.204783
      Name: smoker, dtype: float64
```

## 0.1 SUMMARY FROM EXPLORING DATA

- This dataset contains 1,338 records, with 6 independent features and an individual medical costs charged by health insurance as dependent variable
- This dataset contains no null values in any column
- Out of 6 features, 3 are numerical features: age, bmi and number of children, and
- Out of 6 features, 3 are categorical features: sex, smoker status and region
- Graphing multiple pairplots, we see that the feature with most impact on medical cost/charges is the smoker status
- For smoker, the interquartile cost ranges from 20k to 40k, with the median of around 36k
- For non-smoker, the interquartile cost ranges from only 4k to 11k, with the median of around 7k
- This mean that the medical cost for smokers could be as high as 5 times the cost for non-smokers!
- The slight problem from this summary might be the fact that sample size of smokers accounts for only 20% of entire dataset

```
[29]: X = data.iloc[:, data.columns != 'charges']
      X.head()
```

```
[29]:   age    sex    bmi  children  smoker    region
0   19  female  27.900         0     yes  southwest
1   18   male  33.770         1     no   southeast
2   28   male  33.000         3     no   southeast
3   33   male  22.705         0     no  northwest
4   32   male  28.880         0     no  northwest
```

```
[30]: y= data.iloc[:, data.columns == 'charges']
      y.head()
```

```
[30]:      charges
0  16884.92400
1   1725.55230
2   4449.46200
3  21984.47061
4   3866.85520
```

```
[31]: X_numerical_col = X.select_dtypes(include =[int, float]).columns
X_categorical_col = X.select_dtypes(exclude =[int, float]).columns

print("Numerical column name: {}".format(X_numerical_col))
print("Numerical column name: {}".format(X_categorical_col))
```

```
Numerical column name: Index(['age', 'bmi', 'children'], dtype='object')
Numerical column name: Index(['sex', 'smoker', 'region'], dtype='object')
```

## 0.2 Preprocessing Categorical Columns with OneHotEncoder

```
[32]: from sklearn.preprocessing import OneHotEncoder
encoder = OneHotEncoder(sparse = False)
encoded = encoder.fit_transform(X[X_categorical_col ])
print(encoder.get_feature_names())

#create new encoded dataframe
X_categorical = pd.DataFrame(encoded, columns = encoder.get_feature_names())
X_categorical.head()
```

```
['x0_female' 'x0_male' 'x1_no' 'x1_yes' 'x2_northeast' 'x2_northwest'
 'x2_southeast' 'x2_southwest']
```

```
[32]:
```

	x0_female	x0_male	x1_no	x1_yes	x2_northeast	x2_northwest	\
0	1.0	0.0	0.0	1.0	0.0	0.0	
1	0.0	1.0	1.0	0.0	0.0	0.0	
2	0.0	1.0	1.0	0.0	0.0	0.0	
3	0.0	1.0	1.0	0.0	0.0	1.0	
4	0.0	1.0	1.0	0.0	0.0	1.0	

	x2_southeast	x2_southwest
0	0.0	1.0
1	1.0	0.0
2	1.0	0.0
3	0.0	0.0
4	0.0	0.0

```
[33]: processed_X = X.drop(columns = ['sex', 'smoker', 'region'])
```

```
[34]: processed_X = pd.concat([processed_X, X_categorical], axis = 1)
processed_X.head()
```

```
[34]:
```

	age	bmi	children	x0_female	x0_male	x1_no	x1_yes	x2_northeast	\
0	19	27.900	0	1.0	0.0	0.0	1.0	0.0	
1	18	33.770	1	0.0	1.0	1.0	0.0	0.0	
2	28	33.000	3	0.0	1.0	1.0	0.0	0.0	
3	33	22.705	0	0.0	1.0	1.0	0.0	0.0	

4	32	28.880	0	0.0	1.0	1.0	0.0	0.0
---	----	--------	---	-----	-----	-----	-----	-----

	x2_northwest	x2_southeast	x2_southwest
0	0.0	0.0	1.0
1	0.0	1.0	0.0
2	0.0	1.0	0.0
3	1.0	0.0	0.0
4	1.0	0.0	0.0

### 0.3 Preprocessing with TrainTestSplit

```
[35]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(processed_X, y, test_size = 0.3)
```

```
[36]: print(X_train.shape)
print(y_train.shape)
```

(936, 11)

(936, 1)

### 0.4 Preprocessing Numerical Columns with StandardScaler

The reasons to standardize data AFTER the split is to prevent data leakage from test set!

```
[37]: print(f"The numerical columns to be standardized is {X_numerical_col}")
```

The numerical columns to be standardized is Index(['age', 'bmi', 'children'], dtype='object')

```
[38]: X_mean = X_train[X_numerical_col].mean(axis = 0)
X_std = X_train[X_numerical_col].std(axis = 0)
```

```
[39]: X_train.head()
```

```
[39]:
```

	age	bmi	children	x0_female	x0_male	x1_no	x1_yes	x2_northeast	\
888	22	39.50	0	0.0	1.0	1.0	0.0	0.0	
660	37	46.53	3	0.0	1.0	1.0	0.0	0.0	
934	32	37.18	2	0.0	1.0	1.0	0.0	0.0	
1000	30	22.99	2	0.0	1.0	0.0	1.0	0.0	
97	55	38.28	0	0.0	1.0	1.0	0.0	0.0	

	x2_northwest	x2_southeast	x2_southwest
888	0.0	0.0	1.0
660	0.0	1.0	0.0
934	0.0	1.0	0.0
1000	1.0	0.0	0.0

```
97          0.0          1.0          0.0
```

```
[40]: X_train[X_numerical_col] = (X_train[X_numerical_col] - X_mean) / X_std
```

```
/opt/conda/lib/python3.8/site-packages/pandas/core/frame.py:2963:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
self[k1] = value[k2]
```

```
[41]: X_train.head()
```

```
[41]:
```

	age	bmi	children	x0_female	x0_male	x1_no	x1_yes	\
888	-1.216077	1.400738	-0.920626	0.0	1.0	1.0	0.0	
660	-0.166162	2.551015	1.572255	0.0	1.0	1.0	0.0	
934	-0.516134	1.021130	0.741295	0.0	1.0	1.0	0.0	
1000	-0.656122	-1.300696	0.741295	0.0	1.0	0.0	1.0	
97	1.093737	1.201116	-0.920626	0.0	1.0	1.0	0.0	

	x2_northeast	x2_northwest	x2_southeast	x2_southwest
888	0.0	0.0	0.0	1.0
660	0.0	0.0	1.0	0.0
934	0.0	0.0	1.0	0.0
1000	0.0	1.0	0.0	0.0
97	0.0	0.0	1.0	0.0

```
[42]: y_train.mean()
```

```
[42]: charges    13332.657562  
dtype: float64
```

```
[43]: #scaling on y  
y_mean = y_train.mean()  
y_std = y_train.std()
```

```
[44]: y_train = (y_train - y_mean) / y_std
```

```
[45]: y_train.head()
```

```
[45]:
```

	charges
888	-0.944175
660	-0.558968
934	-0.701787
1000	0.326538



97    -0.251755

## 0.5 Create Linear Regression Model

Convert all the dataframe into numpy

```
[46]: X_train_ = X_train.to_numpy()
      print(X_train_.shape)

      y_train_ = y_train.to_numpy()
      print(y_train_.shape)
```

(936, 11)

(936, 1)

```
[47]: #add intercept on the first index
      X_train_ = np.insert(X_train_, 0,1 ,axis =1)
      print(X_train_.shape)
```

(936, 12)

```
[48]: theta = np.zeros(X_train_.shape[1])
      print(theta.shape)
```

(12,)

```
[49]: #must flatten the y_train so that it will become (m,) instead of (m,1)
      y_train_ = y_train_.flatten()
```

```
[50]: predicted = np.dot(X_train_, theta)
      print(predicted.shape)

      error = predicted - y_train_
      print(error.shape)
```

(936,)

(936,)

```
[51]: #To facilitate the computation, the threshold is set to early stop the gradient_
      ↪descent when the delta between loss < threshold
```

```
theta = theta = np.zeros(X_train_.shape[1])
threshold = 0.00000001
loss_record = [100,10]
count=0
r2_record = []
for i in range(100000):
    alpha = 0.0001
```

```

predicted = np.dot(X_train_, theta)
error = predicted - y_train_

loss = np.sum(error **2) / 2
gradient = np.dot(X_train_.T, error)
theta = theta - (alpha * gradient)
loss_record.append(loss)
r2 = 1 - np.sum((y_train_ - predicted) **2) / np.sum((y_train_ - np.
↪mean(y_train_)) **2)
r2_record.append(r2)
count +=1
if np.abs(loss_record[-1] - loss_record[-2]) < threshold:
    break
else:
    pass

if count % 50 ==0:
    print(f'{count}, loss: {loss}')

print(count)

```

```

50, loss: 128.41181349876595
100, loss: 110.29430994962246
150, loss: 109.08926523294164
200, loss: 109.00529897602448
250, loss: 108.99908625853412
300, loss: 108.99859178493294
350, loss: 108.99854921015694
400, loss: 108.99854526134013
416

```

## 0.6 Plot loss function against iteration for training data

```

[52]: fig,axes = plt.subplots(1,2, figsize = (18,5))
axes[0].plot(np.arange(count), loss_record[2:])
axes[0].set_xlabel('iterations')
axes[0].set_ylabel('loss function')
axes[0].set_title('Loss function against iterations')


axes[1].plot(np.arange(count), r2_record)
axes[1].set_xlabel('iterations')
axes[1].set_ylabel('r2')
axes[1].set_title('r2 function against iterations')

```

```

[52]: Text(0.5, 1.0, 'r2 function against iterations')

```



```
[53]: theta
```

```
[53]: array([ 0.26594193,  0.29631842,  0.1631706 ,  0.04390204,  0.1315308 ,
          0.13441112, -0.85820643,  1.12414835,  0.13653553,  0.05243814,
          0.0451931 ,  0.03177515])
```

```
[54]: r2
```

```
[54]: 0.7668480320081696
```

## 0.7 Summary on performing linear regression on training data

- with alpha at 0.0001 and early stop in place, the model iterates for 416 times at the loss delta < 0.0001
- Using r2 as the metric, the score of this model is around 77%

## 0.8 Perform on the test data

```
[55]: #must standardize the data using the mean and std of training dataset
X_test[X_numerical_col] = (X_test[X_numerical_col] - X_mean) / X_std
```

```
/opt/conda/lib/python3.8/site-packages/pandas/core/frame.py:2963:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
self[k1] = value[k2]
```

```
[56]: X_test.head()
```

```
[56]:      age      bmi  children  x0_female  x0_male  x1_no  x1_yes  \
840 -1.286071  0.026295 -0.920626        0.0        1.0        1.0        0.0
772  0.323799  0.906593 -0.920626        1.0        0.0        1.0        0.0
```

683	0.953748	-1.083076	-0.920626	0.0	1.0	1.0	0.0
1039	-1.426060	-0.601203	0.741295	0.0	1.0	1.0	0.0
987	0.393793	-0.539026	-0.089666	1.0	0.0	1.0	0.0

	x2_northeast	x2_northwest	x2_southeast	x2_southwest
840	0.0	0.0	0.0	1.0
772	1.0	0.0	0.0	0.0
683	0.0	1.0	0.0	0.0
1039	0.0	1.0	0.0	0.0
987	0.0	1.0	0.0	0.0

```
[57]: y_test
```

```
[57]:      charges
840    1526.31200
772   12797.20962
683    9863.47180
1039   22493.65964
987   28340.18885
...
1327   9377.90470
737    3484.33100
767    7050.64200
515   11362.75500
195    1639.56310
```

```
[402 rows x 1 columns]
```

```
[58]: y_test = (y_test - y_mean) / y_std
y_test.head()
```

```
[58]:      charges
840   -0.956841
772   -0.043395
683   -0.281159
1039   0.742450
987    1.216280
```

```
[59]: X_test_ = X_test.to_numpy()
y_test_ = y_test.to_numpy().flatten()
print(X_test_.shape)
print(y_test_.shape)
```

```
(402, 11)
```

```
(402,)
```

```
[60]: X_test_ = np.insert(X_test_, 0,1, axis =1)
      print(X_test_.shape)
```

```
(402, 12)
```

```
[61]: yhat = np.dot(X_test_, theta)
```

```
[62]: r2_yhat = 1 - np.sum(y_test_ - yhat)**2 / np.sum((y_test_ - np.
      ↪mean(y_test_))**2)
      print(r2_yhat)
```

```
0.8030244597975953
```

## 1 r2 of test model is at around 80%

```
[63]: theta
```

```
[63]: array([ 0.26594193,  0.29631842,  0.1631706 ,  0.04390204,  0.1315308 ,
            0.13441112, -0.85820643,  1.12414835,  0.13653553,  0.05243814,
            0.0451931 ,  0.03177515])
```

```
[64]: processed_X.columns
```

```
[64]: Index(['age', 'bmi', 'children', 'x0_female', 'x0_male', 'x1_no', 'x1_yes',
            'x2_northeast', 'x2_northwest', 'x2_southeast', 'x2_southwest'],
            dtype='object')
```

```
[65]: pd.DataFrame(np.array([theta]), columns = ['intercept', 'age', 'bmi',
      ↪'children', 'x0_female', 'x0_male',
            'x1_no', 'x1_yes', 'x2_northeast', 'x2_northwest', 'x2_southeast',
            'x2_southwest'])
```

```
[65]:   intercept      age      bmi  children  x0_female  x0_male  x1_no  \
0   0.265942  0.296318  0.163171  0.043902   0.131531  0.134411 -0.858206

      x1_yes  x2_northeast  x2_northwest  x2_southeast  x2_southwest
0   1.124148      0.136536      0.052438      0.045193      0.031775
```


```
[66]: pd.DataFrame(np.array([theta]), columns = np.insert(processed_X.columns.
      ↪to_numpy(), 0, 'intercept'))
```

```
[66]:   intercept      age      bmi  children  x0_female  x0_male  x1_no  \
0   0.265942  0.296318  0.163171  0.043902   0.131531  0.134411 -0.858206

      x1_yes  x2_northeast  x2_northwest  x2_southeast  x2_southwest
0   1.124148      0.136536      0.052438      0.045193      0.031775
```

```
[67]: sns.heatmap(np.abs(pd.DataFrame(np.array([theta]), columns = np.
      ↳insert(processed_X.columns.to_numpy(), 0, 'intercept'))).T)
```

```
[67]: <matplotlib.axes._subplots.AxesSubplot at 0x7f186c08b760>
```



The heatmap highlights the above-mentioned summary that ‘smoker’ status is the important feature to affect the medical costs

## 1.1 Now try analytical solution

I need to reset the first index ‘intercept’ to something close to 1 but 1 to avoid the multicollinearity that will return an error when trying to find the inverse matrix

I first tried with 1 as normal to find the repeated error of ‘singular matrix’ because in my X\_train\_, it looks like there are some columns that are linearly dependent on other columns

For example, for the ‘intercept’ column, the computer deems it dependent on the sum of One-hot-encoded male/ female or smoke / no smoke!!!

To avoid the singular error, I need to avoid assigning ‘intercept’ as the sum of any number of one-hot-encoded, so I used float number close to 1

The below is the old code that returns error from singular matrix (as reference to see how it blocks me from further coding)

```
[280]: #DO NOT RUN THIS CODE; it just to serves as reference for how the singular
      ↳matrix could not be inverted
```

```
analytical_theta = np.linalg.inv(np.dot(X_train_.T, X_train_)) @ X_train_.T @
↳ y_train_
```

```
↳
-----

LinAlgError                                Traceback (most recent call
↳ last)

  /tmp/ipykernel_167/574514158.py in <module>
    ----> 1 analytical_theta = np.linalg.inv(np.dot(X_train_.T, X_train_)) @
↳ X_train_.T @ y_train_

    <__array_function__ internals> in inv(*args, **kwargs)

  /opt/conda/lib/python3.8/site-packages/numpy/linalg/linalg.py in inv(a)
    543     signature = 'D->D' if isComplexType(t) else 'd->d'
    544     extobj = get_linalg_error_extobj(_raise_linalgerror_singular)
--> 545     ainv = _umath_linalg.inv(a, signature=signature, extobj=extobj)
    546     return wrap(ainv.astype(result_t, copy=False))
    547

  /opt/conda/lib/python3.8/site-packages/numpy/linalg/linalg.py in
↳ _raise_linalgerror_singular(err, flag)
    86
    87 def _raise_linalgerror_singular(err, flag):
--> 88     raise LinAlgError("Singular matrix")
    89
    90 def _raise_linalgerror_nonposdef(err, flag):

LinAlgError: Singular matrix
```

[286]: *#DO NOT RUN THIS CODE; it just serves as reference for how the singular matrix*  
*↳ could not be inverted*  
*#see that intercept columns ~ (X\_female + X-male)*  
X\_train.T @ X\_train

[286]:

	intercept	age	bmi	children	\
intercept	9.360000e+02	1.243450e-14	2.209344e-12	7.549517e-15	
age	1.243450e-14	9.350000e+02	1.067586e+02	2.631458e+01	
bmi	2.209344e-12	1.067586e+02	9.350000e+02	1.090044e+01	

children	7.549517e-15	2.631458e+01	1.090044e+01	9.350000e+02
x0_female	4.470000e+02	2.541935e+01	-1.306959e+01	-2.031536e+01
x0_male	4.890000e+02	-2.541935e+01	1.306959e+01	2.031536e+01
x1_no	7.550000e+02	-1.732732e+00	3.482977e+00	1.824681e+00
x1_yes	1.810000e+02	1.732732e+00	-3.482977e+00	-1.824681e+00
x2_northeast	2.290000e+02	8.038445e-01	-5.544868e+01	-2.538266e+01
x2_northwest	2.500000e+02	-1.082511e+01	-5.513757e+01	1.164486e+01
x2_southeast	2.360000e+02	6.287851e+00	1.099656e+02	-5.703107e+00
x2_southwest	2.210000e+02	3.733411e+00	6.206967e-01	1.944091e+01

	x0_female	x0_male	x1_no	x1_yes	x2_northeast \
intercept	447.000000	489.000000	755.000000	181.000000	229.000000
age	25.419350	-25.419350	-1.732732	1.732732	0.803845
bmi	-13.069588	13.069588	3.482977	-3.482977	-55.448679
children	-20.315360	20.315360	1.824681	-1.824681	-25.382659
x0_female	447.000000	0.000000	377.000000	70.000000	106.000000
x0_male	0.000000	489.000000	378.000000	111.000000	123.000000
x1_no	377.000000	378.000000	755.000000	0.000000	181.000000
x1_yes	70.000000	111.000000	0.000000	181.000000	48.000000
x2_northeast	106.000000	123.000000	181.000000	48.000000	229.000000
x2_northwest	120.000000	130.000000	208.000000	42.000000	0.000000
x2_southeast	112.000000	124.000000	183.000000	53.000000	0.000000
x2_southwest	109.000000	112.000000	183.000000	38.000000	0.000000

	x2_northwest	x2_southeast	x2_southwest
intercept	250.000000	236.000000	221.000000
age	-10.825106	6.287851	3.733411
bmi	-55.137573	109.965555	0.620697
children	11.644860	-5.703107	19.440906
x0_female	120.000000	112.000000	109.000000
x0_male	130.000000	124.000000	112.000000
x1_no	208.000000	183.000000	183.000000
x1_yes	42.000000	53.000000	38.000000
x2_northeast	0.000000	0.000000	0.000000
x2_northwest	250.000000	0.000000	0.000000
x2_southeast	0.000000	236.000000	0.000000
x2_southwest	0.000000	0.000000	221.000000

---

```
[323]: X_train[:,0] = 1.001
```

```
X_train
```

```
[324]: X_train_
```

```
[324]: array([[ 1.001      , -1.21607708,  1.40073753, ...,  0.          ,
              0.          ,  1.          ],
```



```
[ 1.001      , -0.16616181,  2.55101484, ...,  0.      ,
    1.      ,  0.      ],
[ 1.001      , -0.51613356,  1.02112966, ...,  0.      ,
    1.      ,  0.      ],
...,
[ 1.001      , -0.44613921, -0.43676093, ...,  0.      ,
    1.      ,  0.      ],
[ 1.001      , -0.65612227, -1.47168325, ...,  0.      ,
    0.      ,  0.      ],
[ 1.001      ,  0.53378171,  0.22264413, ...,  0.      ,
    0.      ,  1.      ]])
```

```
[325]: np.linalg.inv(np.dot(X_train_.T ,X_train_)) @ X_train_.T @ y_train_
```

```
[325]: array([-1.63208170e+13,  2.94971694e-01,  1.70423816e-01,  4.41698075e-02,
    1.63371378e+13,  1.63371378e+13,  5.43371206e-01,  2.50131628e+00,
   -1.41890509e-01, -2.51389241e-01, -2.48455862e-01, -3.10711286e-01])
```

```
[326]: theta
```

```
[326]: array([ 0.26594193,  0.29631842,  0.1631706 ,  0.04390204,  0.1315308 ,
    0.13441112, -0.85820643,  1.12414835,  0.13653553,  0.05243814,
    0.0451931 ,  0.03177515])
```

```
[327]: analytical_theta = np.linalg.inv(np.dot(X_train_.T ,X_train_)) @ X_train_.T @
↪y_train_
```

### 1.1.1 Analytical method on train data

```
[328]: analytical_yhat_train = np.dot(X_train_, analytical_theta)
analytical_yhat_r2 = 1 - np.sum((y_train_ - analytical_yhat_train)**2) / np.
↪sum((y_train_ - np.mean(y_train_))**2)
```

```
[329]: sse = np.sum((y_train_ - analytical_yhat_train)**2)
tse = np.sum((y_train_ - np.mean(y_train_))**2)
print(f"sse: {sse}")
print(f"tse: {tse}")
print(f"r2: {analytical_yhat_r2}")
```

```
sse: 233.41438426728513
tse: 935.0
r2: 0.7503589473077165
```

### 1.1.2 Analytical method on test data

```
[330]: X_test[:,0] = 1.001
```

```
[331]: analytical_yhat_test = np.dot(X_test_, analytical_theta)
analytical_yhat_r2_test = 1 - np.sum((y_test_ - analytical_yhat_test)**2) / np.
↪sum((y_test_ - np.mean(y_test_))**2)
```

```
[332]: sse = np.sum((y_test_ - analytical_yhat_test)**2)
tse = np.sum((y_test_ - np.mean(y_test_))**2)
print(f"sse: {sse}")
print(f"tse: {tse}")
print(f"r2: {analytical_yhat_r2_test}")
```

sse: 108.09617334970144

tse: 352.78158601985206

r2: 0.6935889580596801

## 1.2 SUMMARY OF LINEAR REGRESSION USING GRADIENT DESCENT vs NORMAL EQUATIONS

- r2 of gradient descent method on training data ~ 77% and ~80% on the test data
  - r2 of analytical method on training data ~ 75% and ~69% on the test data
  - the performance of gradient descent method is slightly better on training data, but noticeably higher than
-