

Before you turn this problem in, make sure everything runs as expected. First, **restart the kernel** (in the menubar, select Kernel→Restart) and then **run all cells** (in the menubar, select Cell→Run All).

Make sure you fill in any place that says YOUR CODE HERE or "YOUR ANSWER HERE", as well as your name and collaborators below:

```
In [1]: NAME = "Lin Tun Naing"
        ID = "st122403"
```

## Lab 05: Optimization Using Newton's Method

In this lab, we'll explore an alternative to gradient descent for nonlinear optimization problems: Newton's method.

### Newton's method in one dimension

Consider the problem of finding the *roots*  $\mathbf{x}$  of a nonlinear function  $f : \mathbb{R}^N \rightarrow \mathbb{R}$ . A root of  $f$  is a point  $\mathbf{x}$  that satisfies  $f(\mathbf{x}) = 0$ .

In one dimension, Newton's method for finding zeroes works as follows:

1. Pick an initial guess  $x_0$
2. Let  $x_{i+1} = x_i + \frac{f(x_i)}{f'(x_i)}$
3. If not converged, go to #2.

Convergence occurs when  $|f(x_i)| < \epsilon_1$  or when  $|f(x_{i+1}) - f(x_i)| < \epsilon_2$ .

Let's see how this works in practice.

### Example 1: Root finding for a cubic polynomial

Let's begin by using Newton's method to find roots of a simple cubic polynomial

$$f(x) = x^3 + x^2.$$

```
In [2]: import matplotlib.pyplot as plt
import numpy as np
from mpl_toolkits.mplot3d import Axes3D
import pandas as pd
```

Here's a function to evaluate a polynomial created with Numpy's `poly1d` function at a particular point  $x$ :

```
In [3]: def fx(x, p):
        f_x = np.polyval(p, x)
        return f_x
```

And here's some code to create the polynomial  $x^3 + x^2$ , get its derivative, and evaluate the derivative at 200 points along the  $x$  axis::

```
In [4]: # Create the polynomial f(x) = x^3 + x^2
p = np.poly1d([1, 1, 0, 0]) # [1 * x^3, 1 * x^2, 0 * x^1, 0 * 1]

# Get f'(x) (the derivative of f(x) in polynomial form)
# We know it's 2x^2 + 2x, which is [3, 2, 0] in poly1d form
p_d = np.polyder(p)

print('f(x):')
print('-----')
print(p)
print('-----')
print("f'(x):")
print('-----')
print(p_d)
print('-----')

# Get 200 points along the x axis between -3 and 3
n = 200
x = np.linspace(-3, 3, n)

# Get values for f(x) and f'(x) in order to graph them later
y = fx(x, p)
y_d = fx(x, p_d)
```

f(x):  
-----  
      3      2  
1 x + 1 x  
-----  
f'(x):  
-----  
      2  
3 x + 2 x  
-----

Next, let's try three possible guesses for  $x_0$ : -3, 1, and 3, and in each case, run Newton's root finding method from that initial guess.

```
In [5]: # Initial guesses
x0_arr = [-3.0, 1.0, 3.0]


# Parameters for Newton: number of iterations,
# threshold for identifying a point as a zero
max_iters = 30
threshold = 0.0001

# Set up plot
fig1 = plt.figure(figsize=(8,8))
ax = plt.axes()
plt.plot(x, y, 'g-', label='f(x)')
plt.plot(x, y_d, 'b--', label="f'(x)")

roots = []
for x0 in x0_arr:
    i = 0
    xi = x0
    fxi = fx(xi, p)
    # Plot initial data point
    plt.plot(xi, fxi, '*', label=("iter 0: x=%.4f" % x0))
    while i < max_iters:
        # x_{i+1} = x_i - f(x_i)/f'(x_i)
        xi = xi - fx(xi, p) / fx(xi, p_d)
        fxi = fx(xi, p)
        # Plot (xi, fxi) and add a Legend entry every 5 iterations
        if (i+1) % 5 == 0:
            plt.plot(xi, fxi, '.', label=("iter %d: x=%.4f" % (i+1, xi)))
        else:
            plt.plot(xi, fxi, '.')
        # Check if |f(x)| < threshold
        if np.abs(fxi) < threshold:
            roots.append(xi)
            break
        i = i + 1
    plt.plot(xi, fx(xi, p), 'ko', label=("converged at iter %d" % (i+1)))

plt.legend(bbox_to_anchor=(1.5, 1.0), loc='upper right')
plt.title('Example 1: Newton root finding for the polynomial x^3 + x^2')

plt.show()
```



**Example 2: Root finding for the sine function**

Next, consider the function  $f(x) = \sin(x)$ :

```
In [6]: def fx_sin(x):
        return np.sin(x)

def fx_dsin(x):
    return np.cos(x)
```

Let's get 200 points in the range  $[-\pi, \pi]$  for plotting:

```
In [7]: # Get f(x)=sin(x) and f'(x) at 200 points for plotting
n = 200
x = np.linspace(-np.pi, np.pi, n)
y = fx_sin(x)
y_d = fx_dsin(x)
```

```
In [8]: # Initial guesses
x0_arr = [2.0, 1.0, -2.0]

# Parameters for Newton: number of iterations,
# threshold for identifying a point as a zero
max_iters = 30
threshold = 0.0001


# Set up plot
fig1 = plt.figure(figsize=(10,10))
ax = plt.axes()
ax.set_aspect(aspect='equal', adjustable='box')
plt.plot(x, y, 'g-', label='f(x)')
plt.plot(x, y_d, 'b--', label='df(x)')

roots = []
for x0 in x0_arr:
    i = 0;
    xi = x0
    fxi = fx_sin(xi)
    # Plot initial data point
    plt.plot(xi, fxi, 'bo', label=("iter 0: x=%.4f" % x0))
    while i < max_iters:
        # x_{i+1} = x_i - f(x_i)/f'(x_i)
        xi = xi - fx_sin(xi) / fx_dsin(xi)
        fxi = fx_sin(xi)
        # Plot (xi, fxi) and add a Legend entry every 5 iterations
        if (i+1) % 5 == 0:
            plt.plot(xi, fxi, '.', label=("iter %d: x=%.4f" % (i+1, xi)))
        else:
            plt.plot(xi, fxi, '.')
        # Check if |f(x)| < threshold
        if np.abs(fxi) < threshold:
            roots.append(xi)
            break
        i = i + 1
    plt.plot(xi, fx_sin(xi), 'ko', label=("converged at iter %d" % (i+1)))

plt.legend(bbox_to_anchor=(1.5, 1.0), loc='upper right')
plt.title('Example 2: Newton root finding for sin(x)')

plt.show()

print('Roots: %f, %f, %f' % (roots[0], roots[1], roots[2]))
```



Roots: 3.141593, -0.000096, -3.141593

Notice that we get some extreme values of  $x$  for some cases. For example, when  $x_0 = -2$ , where the slope is pretty close to 0, the next iteration gives a value less than -4.

## Newton's method for optimization

Now, consider the problem of minimizing a scalar function  $J : \mathbb{R}^n \mapsto \mathbb{R}$ . We would like to find  $\theta^* = \operatorname{argmin}_{\theta} J(\theta)$

We already know gradient descent:

$$\theta^{(i+1)} \leftarrow \theta^{(i)} - \alpha \nabla_J(\theta^{(i)}).$$

But Newton's method gives us a potentially faster way to find  $\theta^*$  as a zero of the system of equations

$$\nabla_J(\theta^*) = \mathbf{0}.$$

In one dimension, to find the zero of  $f'(x)$ , obviously, we would apply Newton's method to  $f'(x)$ , obtaining the iteration

$$x_{i+1} = x_i - f'(x_i)/f''(x_i).$$

The multivariate extension of Newton's optimization method is

$$\mathbf{x}_{i+1} = \mathbf{x}_i - \mathbf{H}_f^{-1}(\mathbf{x}_i) \nabla_f(\mathbf{x}_i),$$

where  $\mathbf{H}_f(\mathbf{x})$  is the *Hessian* of  $f$  evaluated at  $\mathbf{x}$ :

$$\mathbf{H}_f(\mathbf{x}) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 x_n} \\ \frac{\partial^2 f}{\partial x_2 x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n x_1} & \frac{\partial^2 f}{\partial x_n x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

This means, for the minimization of  $J(\theta)$ , we would obtain the update rule

$$\theta^{(i+1)} \leftarrow \theta^{(i)} - \mathbf{H}_J^{-1}(\theta^{(i)}) \nabla_J(\theta^{(i)}).$$

## Application to logistic regression

Let's create some difficult sample data as follows:

**Class 1:** Two features  $x_1$  and  $x_2$  jointly distributed as a two-dimensional spherical Gaussian with parameters

$$\mu = \begin{bmatrix} x_{1c} \\ x_{2c} \end{bmatrix}, \Sigma = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_1^2 \end{bmatrix}.$$

**Class 2:** Two features  $x_1$  and  $x_2$  in which the data are generated by first sampling an angle  $\theta$  according to a uniform distribution, sampling a distance  $d$  according to a one-dimensional Gaussian with a mean of  $(3\sigma_1)^2$  and a variance of  $(\frac{1}{2}\sigma_1)^2$ , then outputting the point

$$\mathbf{x} = \begin{bmatrix} x_{1c} + d \cos \theta \\ x_{2c} + d \sin \theta \end{bmatrix}.$$

Generate 100 samples for each of the classes, guided by the following exercises.

### Exercise 1.1 (5 points)

Generate data for class 1 with 100 samples:

$$\mu = \begin{bmatrix} x_{1c} \\ x_{2c} \end{bmatrix}, \Sigma = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_1^2 \end{bmatrix}.$$

**Hint:**

```
In [9]: mu_1 = np.array([1.0, 2.0])
sigma_1 = 1
num_sample = 100

cov_mat = np.array([[sigma_1, 0], [0, sigma_1]])
X1 = np.random.multivariate_normal(mu_1, cov_mat, num_sample)
```

```
In [10]: print(X1[:5])

# Test function: Do not remove
assert X1.shape == (100, 2), 'Size of X1 is incorrect'
assert cov_mat.shape == (2, 2), 'Size of x_test is incorrect'
count = 0
for i in range(2):
    for j in range(2):
        if i==j and cov_mat[i,j] != 0:
            if cov_mat[i,j] == sigma_1:
                count += 1
        else:
            if cov_mat[i,j] == 0:
                count += 1
assert count == 4, 'cov_mat data is incorrect'

print("success!")
# End Test function
```

```
[[1.20843043 2.10696693]
 [1.14318802 2.00814591]
 [1.47051804 0.68384279]
 [0.51191179 3.27545311]
 [1.02435511 2.10586304]]
success!
```

**Expected result (or something similar):**

```
[[ -0.48508229  2.65415886]
 [ 1.17230227  1.61743589]
 [-0.61932146  3.53986541]
 [ 0.70583088  1.45944356]
 [-0.93561505  0.2042285 ]]
```

**Exercise 1.2 (5 points)**

Generate data for class 2 with 100 samples:

$$\mathbf{x} = \begin{bmatrix} x_{1c} + d \cos \theta \\ x_{2c} + d \sin \theta \end{bmatrix}$$

where  $\theta$  is sampled uniformly from  $[0, 2\pi]$  and  $d$  is sampled from a one-dimensional Gaussian with a mean of  $(3\sigma_1)^2$  and a variance of  $(\frac{1}{2}\sigma_1)^2$ .

**Hint:**

```
In [11]: # 1. Create sample angle from 0 to 2pi with 100 samples
angle = np.random.uniform(0, 2*np.pi, 100)
# 2. Create sample with normal distribution of d with mean and variance
mean = (3 * sigma_1)**2
variance = (sigma_1/2)**2
d = np.random.normal(mean, variance, 100)
# 3 Create X2
X2 = np.array([X1[:,0] +( d* (np.cos(angle))), X1[:,1] + (d * (np.sin(angle)))]).

# YOUR CODE HERE
# raise NotImplementedError()
X2.shape
```

```
Out[11]: (100, 2)
```

```
In [12]: print('angle:',angle[:5])
print('d:', d[:5])
print('X2:', X2[:5])

# Test function: Do not remove
assert angle.shape == (100,) or angle.shape == (100,1) or angle.shape == 100, 'Size of angle is incorrect'
assert d.shape == (100,) or d.shape == (100,1) or d.shape == 100, 'Size of d is incorrect'
assert X2.shape == (100,2), 'Size of X2 is incorrect'
assert angle.min() >= 0 and angle.max() <= 2*np.pi, 'angle generate incorrect'
assert d.min() >= 8 and d.max() <= 10, 'd generate incorrect'
assert X2[:,0].min() >= -13 and X2[:,0].max() <= 13, 'X2 generate incorrect'
assert X2[:,1].min() >= -10 and X2[:,1].max() <= 13.5, 'X2 generate incorrect'

print("success!")
# End Test function
```

angle: [3.55181342 6.05906779 4.62503987 2.31340236 4.51858335]  
d: [8.67847388 8.84772542 8.9436972 8.81147448 9.27273754]  
X2: [[-6.75001479e+00 9.76963726e+00]  
[ 6.90287128e-01 -5.44649620e+00]  
[-7.61524646e-01 -4.95410181e+00]  
[-2.90878884e+00 6.14898900e-03]  
[-9.88422141e+00 -5.20561188e+00]]  
success!

**Expected result (or something similar):**  
angle: [4.77258271 3.19733552 0.71226709 2.11244845 6.06280915]  
d: [9.13908279 8.84218552 9.24427852 8.74831667 8.85727588]  
X2: [[ 0.064701 -6.46837219]  
[-7.65614929 1.12480234]  
[ 6.37750805 9.58147629]  
[-3.80438416 8.95550952]  
[ 7.70745021 -1.73194274]]

**Exercise 1.3 (5 points)**

Combine X1 and X2 into single dataset

```
In [13]: # 1. concatenate X1, X2 together
X = np.concatenate((X1,X2), axis = 0)
# 2. Create y with class 1 as 0 and class 2 as 1
y = np.concatenate((np.zeros((100,)), np.ones((100,))), axis = 0)

# YOUR CODE HERE
# raise NotImplementedError()
# print(y.shape)
```

```
In [14]: print("shape of X:", X.shape)
print("shape of y:", y.shape)

# Test function: Do not remove
assert X.shape == (200, 2), 'Size of X is incorrect'
assert y.shape == (200,) or y.shape == (200,1) or y.shape == 200, 'Size of y is incorrect'
assert y.min() == 0 and y.max() == 1, 'class type setup is incorrect'

print("success!")
# End Test function
```

shape of X: (200, 2)  
shape of y: (200,)  
success!


**Expect result (or looked alike):**  
shape of X: (200, 2)  
shape of y: (200, 1)

Exercise 1.4 (5 points)

Plot the graph between class1 and class2 with **difference color and point style**.


```
In [15]: fig1 = plt.figure(figsize=(8,8))
ax = plt.axes()
plt.title('Sample data for classification problem')
plt.grid(axis='both', alpha=.25)
# plot graph here
# YOUR CODE HERE
# raise NotImplementedError()
markers = ["o", "*"]
colors = ['b', 'g']

for i, c in enumerate(np.unique(y)):
    plt.scatter(X[y == c, 0], X[y == c, 1], c = colors[i], marker = markers[i], 1)
# end plot graph
plt.legend()
plt.axis('equal')
plt.show()
```



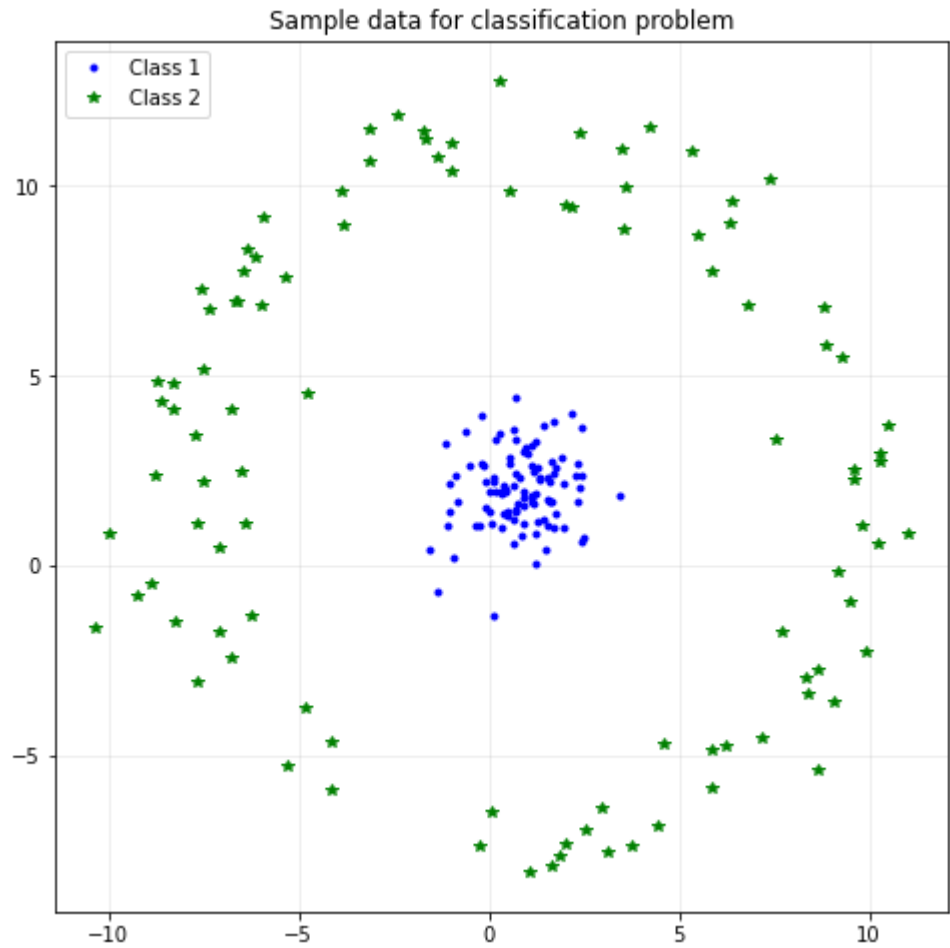
```
In [16]: plt.scatter(X1[:,0], X1[:,1], c='b', marker='o')
plt.scatter(X2[:,0], X2[:,1], c='g', marker='*')
```

Out[16]: <matplotlib.collections.PathCollection at 0x7fe701761160>



Expect result (or looked alike):





Exercise 1.5 (5 points)

Split data into training and test datasets with 80% of training set and 20% of test set

```
In [17]: import random

train_size = 0.8
m = X.shape[0]
idx = np.arange(0, m)
random.shuffle(idx)
train_length = int(train_size * m)

idx_train = idx[:train_length]
idx_test = idx[train_length:]

X_train = X[idx_train, :]
X_test = X[idx_test, :]
y_train = y[idx_train]
y_test = y[idx_test]

# YOUR CODE HERE
# raise NotImplementedError()
# print(y)
```

```
In [18]: print('idx_train:', idx_train[:10])
print("train size, X:", X_train.shape, ", y:", y_train.shape)
print("test size, X:", X_test.shape, ", y:", y_test.shape)

# Test function: Do not remove
assert X_train.shape == (160, 2), 'Size of X_train is incorrect'
assert y_train.shape == (160,) or y_train.shape == (160,1) or y.shape == 160, 'Size of y_train is incorrect'
assert X_test.shape == (40, 2), 'Size of X_test is incorrect'
assert y_test.shape == (40,) or y_test.shape == (40,1) or y.shape == 40, 'Size of y_test is incorrect'

print("success!")
# End Test function
```

```
idx_train: [ 83 117  89  24 104 192  98  41 164 135]
train size, X: (160, 2) , y: (160,)
test size, X: (40, 2) , y: (40,)
success!
```

**Expected result (or something similar):**

idx\_train: [ 78 61 28 166 80 143 6 76 98 133]  
train size, X: (160, 2) , y: (160, 1)  
test size, X: (40, 2) , y: (40, 1)

**Exercise 1.6 (5 points)**

Write a function to normalize your X data

**Practice yourself (No grade, but has extra score 3 points)**

Try to use Jupyter notebook's LaTeX equation capabilities to write the normalization equations for your dataset.

YOUR ANSWER HERE

```
In [19]: def normalization(X):
        """
        Take in numpy array of X values and return normalize X values,
        the mean and standard deviation of each feature
        """

        means = np.mean(X, axis = 0)
        stds = np.std(X, axis = 0)
        X_norm = (X - means) / stds
        # YOUR CODE HERE
        #     raise NotImplementedError()

        return X_norm
```

```
In [20]: XX = normalization(X)

X_train_norm = XX[idx_train]
X_test_norm = XX[idx_test]

# Add 1 at the first column of training dataset (for bias) and use it when training
X_design_train = np.insert(X_train_norm,0,1,axis=1)
X_design_test = np.insert(X_test_norm,0,1,axis=1)

m,n = X_design_train.shape

print(X_train_norm.shape)
print(X_design_train.shape)
print(X_test_norm.shape)
print(X_design_test.shape)

# Test function: Do not remove
assert XX[:,0].min() >= -2.5 and XX[:,0].max() <= 2.5, 'Does the XX is normalized'
assert XX[:,1].min() >= -2.5 and XX[:,1].max() <= 2.5, 'Does the XX is normalized'

print("success!")
# End Test function
```

(160, 2)  
(160, 3)  
(40, 2)  
(40, 3)  
success!

**Exercise 1.7 (10 points)**

define class for logistic regression: batch gradient descent

The class includes:

- **Sigmoid** function

$$sigmoid(z) = \frac{1}{1 + e^{-z}}$$

- **Softmax** function

$$softmax(z) = \frac{e^{z_i}}{\sum_n e^z}$$

- **Hyperthesis (h)** function

$$\hat{y} = h(X; \theta) = softmax(\theta. X)$$

- **Gradient (Negative likelihood)** function

$$gradient = -X. \frac{y - \hat{y}}{n}$$

- **Cost** function

$$cost = \frac{\sum ((-y \log \hat{y}) - ((1 - y) \log (1 - \hat{y})))}{n}$$

- **Gradient ascent** function
- **Prediction** function
- **Get accuracy** funciton

```

In [21]: class Logistic_BGD:
    def __init__(self):
        pass

    def sigmoid(self,z):
        s = 1/(1+ np.exp(-z))
        # YOUR CODE HERE
        # raise NotImplementedError()
        return s

    def softmax(self, z):
        sm = np.exp(z)/(np.sum(np.exp(z)))
        # YOUR CODE HERE
        # raise NotImplementedError()
        return sm

    def h(self,X, theta):
        hf = self.sigmoid(X @ theta)
        # YOUR CODE HERE
        # raise NotImplementedError()
        return hf

    def gradient(self, X, y, y_pred):
        grad = np.dot(X.T, (y_pred - y))/ len(y)
        # YOUR CODE HERE
        # raise NotImplementedError()
        return grad

    def costFunc(self, theta, X, y):
        y_hat = self.h(X, theta)
        cost = (np.sum(-y * np.log(y_hat) - (1 - y) * np.log(1 - y_hat)))/ X.shape[0]
        grad = self.gradient(X, y, y_hat)
        # YOUR CODE HERE
        # raise NotImplementedError()
        return cost, grad

    def gradientAscent(self, X, y, theta, alpha, num_iters):
        m = len(y)
        J_history = []
        theta_history = []
        for i in range(num_iters):
            # 1. calculate cost, grad function
            cost, grad = self.costFunc(theta, X, y)
            # 2. update new theta
            theta = theta - alpha * grad
            # YOUR CODE HERE
            # raise NotImplementedError()

            J_history.append(cost)
            theta_history.append(theta)
        J_min_index = np.argmin(J_history)
        print("Minimum at iteration:",J_min_index)
        return theta_history[J_min_index] , J_history

    def predict(self,X, theta):
        labels=[]
        # 1. take y_predict from hyperthesis function
        # 2. classify y_predict that what it should be class1 or class2
        # 3. append the output from prediction
        # YOUR CODE HERE
        y_predict = self.h(X, theta)
        for i in range(X.shape[0]):
            if y_predict[i] >= 0.5:
                labels.append(1)
            else:
                labels.append(0)

        labels=np.asarray(labels)
        return labels

    def getAccuracy(self,X,y,theta):
        predict_y = self.predict(X, theta)

```

```
correct = 0
for i in range(len(predict_y)):
    if predict_y[i] == y[i]:
        correct += 1
percent_correct = correct / len(y) * 100
# YOUR CODE HERE
# raise NotImplementedError()
return percent_correct
```

```
In [22]: # Test function: Do not remove
lbgd = Logistic_BGD()
test_x = np.array([[1,2,3,4,5]]).T
out_x1 = lbgd.sigmoid(test_x)
out_x2 = lbgd.sigmoid(test_x.T)
print('out_x1', out_x1.T)
assert np.array_equal(np.round(out_x1.T, 5), np.round([[0.73105858, 0.88079708, 0.95257413, 0.98201379, 0.99330715]]))
assert np.array_equal(np.round(out_x2, 5), np.round([[0.73105858, 0.88079708, 0.95257413, 0.98201379, 0.99330715]]))
out_x1 = lbgd.softmax(out_x1)
out_x2 = lbgd.softmax(out_x2)
print('out_x1', out_x1.T)
assert np.array_equal(np.round(out_x1.T, 5), np.round([[0.16681682, 0.19376282, 0.20818183, 0.21440174, 0.21683678]]))
assert np.array_equal(np.round(out_x2, 5), np.round([[0.16681682, 0.19376282, 0.20818183, 0.21440174, 0.21683678]]))
test_t = np.array([[0.3, 0.2]]).T
test_x = np.array([[1,2,3,4,5, 6], [2, 9, 4, 3, 1, 0]]).T
test_y = np.array([[0,1,0,1,0,1]]).T
test_y_p = lbgd.h(test_x, test_t)
print('test_y_p', test_y_p.T)
assert np.array_equal(np.round(test_y_p.T, 5), np.round([[0.66818777, 0.9168273, 0.84553473, 0.85814894, 0.84553473, 0.85814894]]))
test_g = lbgd.gradient(test_x, test_y, test_y_p)
print('test_g', test_g.T)
assert np.array_equal(np.round(test_g.T, 5), np.round([[0.9746016, 0.73165696]])),
test_c, test_g = lbgd.costFunc(test_t, test_x, test_y)
print('test_c', test_c.T)
assert np.round(test_c, 5) == np.round(0.87192491, 5), "costFunc function is incorrect"
test_t_out, test_j = lbgd.gradientAscent(test_x, test_y, test_t, 0.001, 3)
print('test_t_out', test_t_out.T)
print('test_j', test_j)
assert np.array_equal(np.round(test_t_out.T, 5), np.round([[0.29708373, 0.19781153]]))
assert np.round(test_j[2], 5) == np.round(0.86896665, 5), "gradientAscent function is incorrect"
test_l = lbgd.predict(test_x, test_t)
print('test_l', test_l)
assert np.array_equal(np.round(test_l, 1), np.round([1,1,1,1,1,1], 1)), "gradient function is incorrect"
test_a = lbgd.getAccuracy(test_x, test_y, test_t)
print('test_a', test_a)
assert np.round(test_a, 1) == 50.0, "getAccuracy function is incorrect"

print("success!")
# End Test function
```

```
out_x1 [[0.73105858 0.88079708 0.95257413 0.98201379 0.99330715]]
out_x1 [[0.16681682 0.19376282 0.20818183 0.21440174 0.21683678]]
test_y_p [[0.66818777 0.9168273 0.84553473 0.85814894 0.84553473 0.85814894]]
test_g [[0.9746016 0.73165696]]
test_c 0.8719249134773479
Minimum at iteration: 2
test_t_out [[0.29708373 0.19781153]]
test_j [0.8719249134773479, 0.870441756946089, 0.8689666485816598]
test_l [1 1 1 1 1 1]
test_a 50.0
success!
```

**Expected result:**

```
out_x1 [[0.73105858 0.88079708 0.95257413 0.98201379 0.99330715]]
out_x1 [[0.16681682 0.19376282 0.20818183 0.21440174 0.21683678]]
test_y_p [[0.66818777 0.9168273 0.84553473 0.85814894 0.84553473 0.85814894]]
test_g [[0.9746016 0.73165696]]
test_c [0.87192491]
Minimum at iteration: 2
test_t_out [[0.29708373 0.19781153]]
```

```
test_j [array([0.87192491]), array([0.87044176]), array([0.86896665])]
test_l [1 1 1 1 1 1]
test_a 50.0
```

Exercise 1.8 (5 points)

Training the data using Logistic\_BGD class.

- Input: X\_design\_train
- Output: y\_train
- Use 50,000 iterations

Find the initial\_theta yourself

```
In [23]: alpha = 0.001
iterations = 50000

BGD_model = Logistic_BGD()
initial_theta = np.random.randn(X_design_train.shape[1],)
bgd_theta, bgd_cost = BGD_model.gradientAscent(X_design_train, y_train, initial_theta)

# YOUR CODE HERE
# raise NotImplementedError()
```

Minimum at iteration: 49999


```
In [24]: print(bgd_theta)
print(len(bgd_cost))

print(bgd_cost[0])
plt.plot(bgd_cost)
plt.show()

# Test function: Do not remove
assert bgd_theta.shape == (X_train.shape[1] + 1,1) or bgd_theta.shape == (X_train.shape[1],1)
assert len(bgd_cost) == iterations, "cost data size is incorrect"

print("success!")
# End Test function
```


```
[ 0.051114452 -0.01338292 -0.01763408]
50000
1.499471045823371
```



success!

Expected result (or look alike):

```
[[ -0.07328673]
[ -0.13632896]
[ 0.05430939]]
50000
```



**In lab exercises**

1. Verify that the gradient descent solution is correct. Plot the optimal decision boundary you obtain.
2. Write a new class that uses Newton's method for the optimization rather than simple gradient descent.
3. Verify that you obtain a similar solution with Newton's method. Plot the optimal decision boundary you obtain.
4. Compare the number of iterations required for gradient descent vs. Newton's method. Do you observe other issues with Newton's method such as a singular or nearly singular Hessian matrix?

**Exercise 1.9 (5 points)**

Plot the optimal decision boundary of gradient ascent

```
In [25]: # YOUR CODE HERE
def boundary_points(X, theta):
    theta = theta.reshape(-1,1)
    v_orthogonal = np.array([[theta[1,0]], [theta[2,0]]])
    v_ortho_length = np.sqrt(v_orthogonal.T @ v_orthogonal)
    dist_ortho = theta[0,0] / v_ortho_length
    v_orthogonal = v_orthogonal / v_ortho_length
    v_parallel = np.array([[-v_orthogonal[1,0]], [v_orthogonal[0,0]]])
    projections = X @ v_parallel
    proj_1 = min(projections)
    proj_2 = max(projections)
    point_1 = proj_1 * v_parallel - dist_ortho * v_orthogonal
    point_2 = proj_2 * v_parallel - dist_ortho * v_orthogonal
    return point_1, point_2

x_df = pd.DataFrame(X, columns=['X0', 'X1'])

x_df['y'] = y

x_df['X0'] = normalization(x_df.X0)
x_df['X1'] = normalization(x_df.X1)
linX = x_df[['X0', 'X1']].values
linX = np.insert(linX, 0, 1, axis=1)

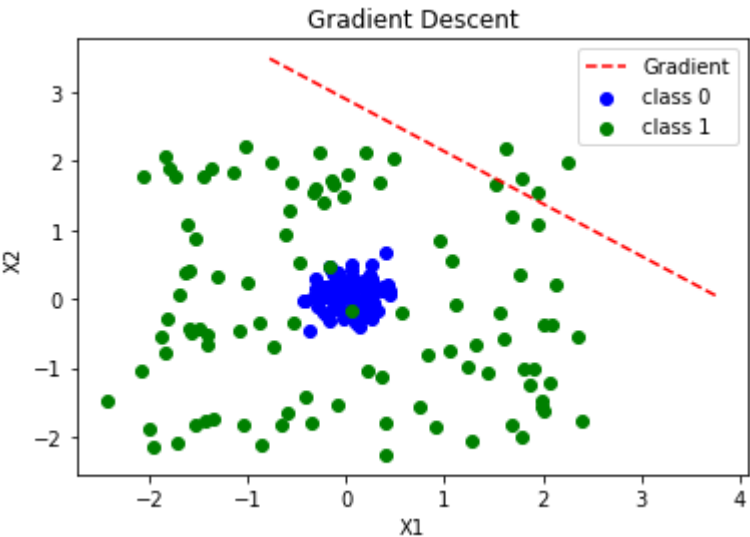
X_train = linX[idx_train]
X_test = linX[idx_test]
y_train = y[idx_train]
y_test = y[idx_test]

y0_df = x_df[x_df.y == 0]
y1_df = x_df[x_df.y == 1]

point_1, point_2 = boundary_points(linX[:,1:], bgd_theta)


plt.title('Gradient Descent')
plt.scatter(y0_df.X0, y0_df.X1, c='b', label='class 0')
plt.scatter(y1_df.X0, y1_df.X1, c='g', label='class 1')
plt.legend()
plt.xlabel('X1')
plt.ylabel('X2')

# plot the boundaries for both methods
plt.plot([point_1[0,0], point_2[0,0]], [point_1[1,0], point_2[1,0]], 'r--', label='Gradient Descent')
plt.legend(loc=0)
plt.show()
```



Expected result (or look alike):





```
In [26]: print("Accuracy =",BGD_model.getAccuracy(X_design_test,y_test,bgd_theta))
```

Accuracy = 45.0

**Exercise 2.1 (10 points)**

Write Newton's method class

In [27]: `class Logistic_NM: #logistic regression for newton's method`

```

    def __init__(self):
        pass

    def sigmoid(self,z):
        s = 1/(1+ np.exp(-z))
        # YOUR CODE HERE
        # raise NotImplementedError()
        return s

    def h(self,X, theta):
        hf = self.sigmoid(X @ theta)
        # YOUR CODE HERE
        # raise NotImplementedError()
        return hf

    def gradient(self, X, y, y_pred):
        grad = np.dot(X.T, (y_pred - y))/ len(y)
        # YOUR CODE HERE
        # raise NotImplementedError()
        return grad

    def hessian(self, X, y, theta):
        y_hat = self.h(X, theta)
        # YOUR CODE HERE
        hess_mat = ((y_hat).T @ (1-y_hat)) * (X.T @ X)/ len(y_hat)
        # raise NotImplementedError()
        return hess_mat

    def costFunc(self, theta, X, y):
        y_hat = self.h(X, theta)
        cost = (np.sum(-y * np.log(y_hat) - (1 - y) * np.log(1 - y_hat)))/ X.shape[0]
        grad = self.gradient(X, y, y_hat)
        # YOUR CODE HERE
        # raise NotImplementedError()
        return cost, grad

    def newtonsMethod(self, X, y, theta, num_iters):
        m = len(y)
        J_history = []
        theta_history = []

        for i in range(num_iters):
            hessian_mat = np.zeros((X.shape[1], X.shape[1]))
            hmat_xi = self.hessian(X, y, theta)
            hessian_mat += hmat_xi
            cost, grad = self.costFunc(theta, X,y)

            try :
                theta = theta - np.linalg.inv(hessian_mat) @ grad
            except Exception as e :
                error_msg = e
                found_singular_matrix = True

            theta = theta - np.linalg.pinv(hessian_mat) @ grad

            J_history.append(cost)
            theta_history.append(theta)
        J_min_index = np.argmin(J_history)
        print("Minimum at iteration:", J_min_index)
        return theta_history[J_min_index] , J_history

    def predict(self,X, theta):
        labels=[]
        # 1. take y_predict from hyperthesis function
        # 2. classify y_predict that what it should be class1 or class2
        # 3. append the output from prediction
        # YOUR CODE HERE
        y_predict = self.h(X, theta)
        for i in range(X.shape[0]):
            if y_predict[i] >= 0.5:
                labels.append(1)

```

```
        else:
            labels.append(0)

        labels=np.asarray(labels)
        return labels

    def getAccuracy(self,X,y,theta):
        predict_y = self.predict(X, theta)
        correct = 0
        for i in range(len(predict_y)):
            if predict_y[i] == y[i]:
                correct += 1
        percent_correct = correct/ len(y) * 100
        # YOUR CODE HERE
        # raise NotImplementedError()
        return percent_correct
```

```
In [28]: # Test function: Do not remove
lbgd = Logistic_NM()
test_x = np.array([[1,2,3,4,5]]).T
out_x1 = lbgd.sigmoid(test_x)
out_x2 = lbgd.sigmoid(test_x.T)
print('out_x1', out_x1.T)
assert np.array_equal(np.round(out_x1.T, 5), np.round([[0.73105858, 0.88079708, 0.95257413, 0.98201379, 0.99330715]], 5))
assert np.array_equal(np.round(out_x2, 5), np.round([[0.73105858, 0.88079708, 0.95257413, 0.98201379, 0.99330715]], 5))
test_t = np.array([[0.3, 0.2]]).T
test_x = np.array([[1,2,3,4,5, 6], [2, 9, 4, 3, 1, 0]]).T
test_y = np.array([[0,1,0,1,0,1]]).T
test_y_p = lbgd.h(test_x, test_t)
print('test_y_p', test_y_p.T)
assert np.array_equal(np.round(test_y_p.T, 5), np.round([[0.66818777, 0.9168273, 0.84553473, 0.85814894, 0.84553473, 0.85814894]], 5))
test_g = lbgd.gradient(test_x, test_y, test_y_p)
print('test_g', test_g.T)
assert np.array_equal(np.round(test_g.T, 5), np.round([[0.9746016, 0.73165696]], 5))
test_h = lbgd.hessian(test_x, test_y, test_t)
print('test_h', test_h)
assert test_h.shape == (2, 2), "hessian matrix function is incorrect"
assert np.array_equal(np.round(test_h.T, 5), np.round([[12.17334371, 6.55487738], [6.55487738, 14.84880387]], 5))
test_c, test_g = lbgd.costFunc(test_t, test_x, test_y)
print('test_c', test_c.T)
assert np.round(test_c, 5) == np.round(0.87192491, 5), "costFunc function is incorrect"
test_t_out , test_j = lbgd.newtonsMethod(test_x, test_y, test_t, 3)
print('test_t_out', test_t_out.T)
print('test_j', test_j)
assert np.array_equal(np.round(test_t_out.T, 5), np.round([[0.14765747, 0.15607017]], 5))
assert np.round(test_j[2], 5) == np.round(0.7534506190845247, 5), "newtonsMethod function is incorrect"
test_l = lbgd.predict(test_x, test_t)
print('test_l', test_l)
assert np.array_equal(np.round(test_l, 1), np.round([1,1,1,1,1,1], 1)), "gradient function is incorrect"
test_a = lbgd.getAccuracy(test_x,test_y,test_t)
print('test_a', test_a)
assert np.round(test_a, 1) == 50.0, "getAccuracy function is incorrect"

print("success!")
# End Test function
```

```
out_x1 [[0.73105858 0.88079708 0.95257413 0.98201379 0.99330715]]
test_y_p [[0.66818777 0.9168273 0.84553473 0.85814894 0.84553473 0.85814894]]
test_g [[0.9746016 0.73165696]]
test_h [[12.17334371 6.55487738]
 [ 6.55487738 14.84880387]]
test_c 0.8719249134773479
Minimum at iteration: 2
test_t_out [[0.14765747 0.15607017]]
test_j [0.8719249134773479, 0.7967484437157274, 0.7534506190845246]
test_l [1 1 1 1 1 1]
test_a 50.0
success!
```

```
Expect result: out_x1 [[0.73105858 0.88079708 0.95257413 0.98201379 0.99330715]]
test_y_p [[0.66818777 0.9168273 0.84553473 0.85814894 0.84553473 0.85814894]]
test_g [[0.9746016 0.73165696]]
```

```
test_h [[12.17334371 6.55487738]
[ 6.55487738 14.84880387]]
test_c 0.8719249134773479
Minimum at iteration: 2
test_t_out [[0.14765747 0.15607017]]
test_j [0.8719249134773479, 0.7967484437157274, 0.7534506190845247]
test_l [1 1 1 1 1]
test_a 50.0
```


```
In [29]: NM_model = Logistic_NM()

iterations = 1000

nm_theta, nm_cost = NM_model.newtonsMethod(X_design_train, y_train, initial_theta)
print("theta:",nm_theta)


print(nm_cost[0])
plt.plot(nm_cost)
plt.show()
```

Minimum at iteration: 999  
theta: [ 0.04893522 -0.01596414 -0.02354655]  
1.499471045823371



**Expected result (or look alike):**

Minimum at iteration: 999  
theta: [[-0.07313861]  
[-0.13605172]  
[ 0.05419746]]  
0.6931471805599453



**Exercise 2.2 (5 points)**

Plot the optimal decision boundary of Newton method

```
In [30]: # YOUR CODE HERE
# raise NotImplementedError()
x_df = pd.DataFrame(X, columns=['X0', 'X1'])

x_df['y'] = y


x_df['X0'] = normalization(x_df.X0)
x_df['X1'] = normalization(x_df.X1)
linX = x_df[['X0', 'X1']].values
linX = np.insert(linX, 0, 1, axis=1)

X_train = linX[idx_train]
X_test = linX[idx_test]
y_train = y[idx_train]
y_test = y[idx_test]


y0_df = x_df[x_df.y == 0]
y1_df = x_df[x_df.y == 1]

point_1n, point_2n = boundary_points(linX[:,1:], nm_theta)

plt.title("Newton's Method")
plt.scatter(y0_df.X0, y0_df.X1, c='g', label='class 0')
plt.scatter(y1_df.X0, y1_df.X1, c='b', label='class 1')
plt.legend()
plt.xlabel('X1')
plt.ylabel('X2')
# plot the boundaries for both methods
plt.plot([point_1n[0,0], point_2n[0,0]],[point_1n[1,0], point_2n[1,0]], 'r--', label='Newton')
plt.legend(loc=0)
plt.show()
```



Expected result (or look alike):



```
In [31]: print("Accuracy =",NM_model.getAccuracy(X_design_test,y_test,nm_theta))
```

Accuracy = 45.0

Exercise 2.3 (5 points)

Compare the number of iterations required for gradient descent vs. Newton's method. Do you observe other issues with Newton's method such as a singular or nearly singular Hessian matrix?

Gradient Descent requires 10000 iterations to reach minimal cost where Newton's Method only requires around 300 iterations to reach the minimum. Sometimes singular or nearly singular Hessian Matrix can be observed.

Take-home exercises

- 1. Perform a *polar transformation* on the data above to obtain a linearly separable dataset. (5 points)
- 2. Verify that you obtain good classification accuracy for logistic regression with GD or Netwon's method after the polar transformation (10 points)
- 3. Apply Newton's method to the dataset you used for the take home exercises in Lab 03. (20 points)

In [32]: `class Logistic_NM: #logistic regression for newton's method`

```

    def __init__(self):
        pass

    def sigmoid(self,z):
        s = 1/(1+ np.exp(-z))
        # YOUR CODE HERE
        # raise NotImplementedError()
        return s

    def h(self,X, theta):
        hf = self.sigmoid(X @ theta)
        # YOUR CODE HERE
        # raise NotImplementedError()
        return hf

    def gradient(self, X, y, y_pred):
        grad = np.dot(X.T, (y_pred - y))/ len(y)
        # YOUR CODE HERE
        # raise NotImplementedError()
        return grad

    def hessian(self, X, y, theta):
        y_hat = self.h(X, theta)
        # YOUR CODE HERE
        hess_mat = ((y_hat).T @ (1-y_hat)) * (X.T @ X)/ len(y_hat)
        # raise NotImplementedError()
        return hess_mat

    def costFunc(self, theta, X, y):
        y_hat = self.h(X, theta)
        cost = (np.sum(-y * np.log(y_hat) - (1 - y) * np.log(1 - y_hat)))/ X.shape[0]
        grad = self.gradient(X, y, y_hat)
        # YOUR CODE HERE
        # raise NotImplementedError()
        return cost, grad

    def newtonsMethod(self, X, y, theta, num_iters):
        m = len(y)
        J_history = []
        theta_history = []

        for i in range(num_iters):
            hessian_mat = np.zeros((X.shape[1], X.shape[1]))
            hmat_xi = self.hessian(X, y, theta)
            hessian_mat += hmat_xi
            cost, grad = self.costFunc(theta, X,y)

            #         try :
            #             theta = theta - np.linalg.inv(hessian_mat) @ grad
            #         except Exception as e :
            #             error_msg = e
            #             found_singular_matrix = True

            theta = theta - np.linalg.pinv(hessian_mat) @ grad

            J_history.append(cost)
            theta_history.append(theta)
        J_min_index = np.argmin(J_history)
        print("Minimum at iteration:", J_min_index)
        return theta_history[J_min_index] , J_history

    def predict(self,X, theta):
        labels=[]
        # 1. take y_predict from hyperthesis function
        # 2. classify y_predict that what it should be class1 or class2
        # 3. append the output from prediction
        # YOUR CODE HERE
        y_predict = self.h(X, theta)
        for i in range(X.shape[0]):
            if y_predict[i] >= 0.5:
                labels.append(1)

```

```
        else:
            labels.append(0)

        labels=np.asarray(labels)
        return labels

    def getAccuracy(self,X,y,theta):
        predict_y = self.predict(X, theta)
        correct = 0
        for i in range(len(predict_y)):
            if predict_y[i] == y[i]:
                correct += 1
        percent_correct = correct/ len(y) * 100
        # YOUR CODE HERE
        # raise NotImplementedError()
        return percent_correct
```


```
In [33]: # Prepare data
# Convert X to angle and radius
df = pd.DataFrame(X, columns=['X0', 'X1'])
df['angles'] = np.arctan(df.X1 / df.X0)
df['radius'] = np.sqrt(df.X0 ** 2 + df.X1 ** 2)
df['y'] = y

newX = df[['angles', 'radius']].values
newX = np.insert(newX, 0, 1, axis=1)
X_train = newX[idx_train]
X_test = newX[idx_test]
y_train = y[idx_train]
y_test = y[idx_test]
```

```
In [34]: y0_df = df[df.y == 0]
y1_df = df[df.y == 1]

plt.title('Transformed Data')
plt.scatter(y0_df.angles, y0_df.radius, c='r', label='class 0')
plt.scatter(y1_df.angles, y1_df.radius, c='b', label='class 1')
plt.legend()
plt.xlabel('angle')
plt.ylabel('radius')
```

Out[34]: Text(0, 0.5, 'radius')



```
In [35]: alpha = 0.01
iterations = 10000
init_theta = np.ones(newX.shape[1])

lg = Logistic_BGD()
theta, j_hist = lg.gradientAscent(X_train, y_train, init_theta, alpha, iterations)
```

Minimum at iteration: 9999




```
In [36]: iterations = 800
init_theta = np.ones(newX.shape[1])

ln = Logistic_NM()
theta_n, j_hist_n = ln.newtonsMethod(X_train, y_train, init_theta, iterations)
# print(theta_n)
```

Minimum at iteration: 799

```
In [37]: plt.plot(j_hist_n, label='Newton\'s Method')
plt.plot(j_hist[:iterations], label='Batch Gradient Descent')
plt.title('Cost History')
plt.xlabel('iteration')
plt.ylabel('cost')
plt.legend()
plt.show()

print('Minimum Cost for each method from polar transformation')
print('Gradient Descent:', np.min(j_hist))
print('Newton\'s Method :', np.min(j_hist_n))
```




Minimum Cost for each method from polar transformation  
Gradient Descent: 0.1010223915711268  
Newton's Method : 0.0914865419900052

```
In [38]: y0_df = df[df.y == 0]
y1_df = df[df.y == 1]

point_1, point_2 = boundary_points(newX[:,1:], theta)
point_1n, point_2n = boundary_points(newX[:,1:], theta_n)

plt.title('Decision Boundary')
plt.scatter(y0_df.angles, y0_df.radius, c='r', label='class 0')
plt.scatter(y1_df.angles, y1_df.radius, c='b', label='class 1')
plt.legend()
plt.xlabel('angle')
plt.ylabel('radius')
# plot the boundaries for both methods
plt.plot([point_1[0,0], point_2[0,0]], [point_1[1,0], point_2[1,0]], 'g-', label='Gradient')
plt.plot([point_1n[0,0], point_2n[0,0]], [point_1n[1,0], point_2n[1,0]], 'c-', label='Newton')
plt.legend(loc=0)
# plt.ylim(0,14)
plt.xlim(-2,2)
plt.show()
```



```
In [39]: print(X_train.shape)
print(theta.shape)
```

(160, 3)  
(3,)

```
In [40]: yg_pred = lg.predict(X_train, theta)
yn_pred = ln.predict(X_train, theta_n)
g_acc = lg.getAccuracy(X_train, y_train, theta)
n_acc = lg.getAccuracy(X_train, y_train, theta_n)
print(n_acc)
print("Train accuracy for polar transformation")
print('Gradient Accuracy : ', g_acc)
print('Newton Accuracy : ', n_acc)

yg_pred = lg.predict(X_test, theta)
yn_pred = ln.predict(X_test, theta_n)
g_acc = lg.getAccuracy(X_test, y_test, theta)
n_acc = lg.getAccuracy(X_test, y_test, theta_n)

print("Test accuracy for polar transformation")
print('Gradient Accuracy : ', g_acc)
print('Newton Accuracy : ', n_acc)
```

97.5  
Train accuracy for polar transformation  
Gradient Accuracy : 97.5  
Newton Accuracy : 97.5  
Test accuracy for polar transformation  
Gradient Accuracy : 97.5  
Newton Accuracy : 97.5

In [41]: *# Import Pandas. You may need to run "pip3 install pandas" at the console if it's*

```
import pandas as pd

# Import the data

data_train = pd.read_csv('train_LoanPrediction.csv')
data_test = pd.read_csv('test_LoanPrediction.csv')

# Start to explore the data

print('Training data shape', data_train.shape)
print('Test data shape', data_test.shape)

print('Training data:\n', data_train)
```

Training data shape (614, 13)

Test data shape (367, 12)

Training data:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
..	...	...	...	...	...	...	
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	
..	...	...	...	...	
609	2900	0.0	71.0	360.0	
610	4106	0.0	40.0	180.0	
611	8072	240.0	253.0	360.0	
612	7583	0.0	187.0	360.0	
613	4583	0.0	133.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y
..	...	...	...
609	1.0	Rural	Y
610	1.0	Rural	Y
611	1.0	Urban	Y
612	1.0	Urban	Y
613	0.0	Semiurban	N

[614 rows x 13 columns]

In [42]: *# Check for missing values in the training and test data*

```
print('Missing values for train data:\n-----\n', data_train.isna().sum())
print('Missing values for test data \n -----\n', data_test.isna().sum())
```

Missing values for train data:  
-----  
Loan\_ID 0  
Gender 13  
Married 3  
Dependents 15  
Education 0  
Self\_Employed 32  
ApplicantIncome 0  
CoapplicantIncome 0  
LoanAmount 22  
Loan\_Amount\_Term 14  
Credit\_History 50  
Property\_Area 0  
Loan\_Status 0  
dtype: int64

Missing values for test data  
-----  
Loan\_ID 0  
Gender 11  
Married 0  
Dependents 10  
Education 0  
Self\_Employed 23  
ApplicantIncome 0  
CoapplicantIncome 0  
LoanAmount 5  
Loan\_Amount\_Term 6  
Credit\_History 29  
Property\_Area 0  
dtype: int64

In [43]: *# Compute ratio of each category value*  
*# Divide the missing values based on ratio*  
*# Fillin the missing values*  
*# Print the values before and after filling the missing values for confirmation*

```
print(data_train['Married'].value_counts())

married = data_train['Married'].value_counts()
print('Elements in Married variable', married.shape)
print('Married ratio ', married[0]/sum(married.values))

def fill_marital_status(data, yes_num_train, no_num_train):
    data['Married'].fillna('Yes', inplace = True, limit = yes_num_train)
    data['Married'].fillna('No', inplace = True, limit = no_num_train)

fill_marital_status(data_train, 2, 1)
print(data_train['Married'].value_counts())
```

Yes 398  
No 213  
Name: Married, dtype: int64  
Elements in Married variable (2,)  
Married ratio 0.6513911620294599  
Yes 400  
No 214  
Name: Married, dtype: int64

```
In [44]: print(data_train['Dependents'].value_counts())
dependent = data_train['Dependents'].value_counts()

print('Dependent ratio 1 ', dependent['0'] / sum(dependent.values))
print('Dependent ratio 2 ', dependent['1'] / sum(dependent.values))
print('Dependent ratio 3 ', dependent['2'] / sum(dependent.values))
print('Dependent ratio 3+ ', dependent['3+'] / sum(dependent.values))

def fill_dependent_status(num_0_train, num_1_train, num_2_train, num_3_train, num_0_test, num_1_test, num_2_test, num_3_test):
    data_train['Dependents'].fillna('0', inplace=True, limit = num_0_train)
    data_train['Dependents'].fillna('1', inplace=True, limit = num_1_train)
    data_train['Dependents'].fillna('2', inplace=True, limit = num_2_train)
    data_train['Dependents'].fillna('3+', inplace=True, limit = num_3_train)
    data_test['Dependents'].fillna('0', inplace=True, limit = num_0_test)
    data_test['Dependents'].fillna('1', inplace=True, limit = num_1_test)
    data_test['Dependents'].fillna('2', inplace=True, limit = num_2_test)
    data_test['Dependents'].fillna('3+', inplace=True, limit = num_3_test)

fill_dependent_status(9, 2, 2, 2, 5, 2, 2, 1)

print(data_train['Dependents'].value_counts())

# Convert category value "3+" to "4"

data_train['Dependents'].replace('3+', 4, inplace = True)
data_test['Dependents'].replace('3+', 4, inplace = True)
```

```
0      345
1      102
2      101
3+       51
Name: Dependents, dtype: int64
Dependent ratio 1   0.5759599332220368
Dependent ratio 2   0.17028380634390652
Dependent ratio 3   0.1686143572621035
Dependent ratio 3+  0.08514190317195326
0      354
1      104
2      103
3+       53
Name: Dependents, dtype: int64
```

```
In [45]: print(data_train['LoanAmount'].value_counts())

LoanAmt = data_train['LoanAmount'].value_counts()

print('mean loan amount ', np.mean(data_train["LoanAmount"]))

loan_amount_mean = np.mean(data_train["LoanAmount"])

data_train['LoanAmount'].fillna(loan_amount_mean, inplace=True, limit = 22)
data_test['LoanAmount'].fillna(loan_amount_mean, inplace=True, limit = 5)
```

```
120.0    20
110.0    17
100.0    15
187.0    12
160.0    12
..
570.0     1
300.0     1
376.0     1
117.0     1
311.0     1
Name: LoanAmount, Length: 203, dtype: int64
mean loan amount   146.41216216216216
```

```
In [46]: print(data_train['Gender'].value_counts())
gender_sum = data_train['Gender'].value_counts()

male_ratio = gender_sum['Male']/sum(gender_sum.values)
female_ratio = gender_sum['Female']/ sum(gender_sum.values)
# print(male_ratio)
# print(female_ratio)

def fill_gender(male_train, female_train, male_test, female_test):
    data_train['Gender'].fillna('Male', inplace= True, limit= male_train)
    data_train['Gender'].fillna('Female', inplace= True, limit= female_train)
    data_test['Gender'].fillna('Male', inplace = True, limit = male_test)
    data_test['Gender'].fillna('Female', inplace = True, limit = female_test)

print(data_train.isnull().sum()['Gender'])
print(data_test.isnull().sum()['Gender'])

train_na = data_train.isnull().sum()['Gender']
test_na = data_test.isnull().sum()['Gender']
male_train = int((male_ratio * train_na).round())
female_train = int((female_ratio * train_na).round())
male_test = int((male_ratio * test_na).round())
female_test = int((female_ratio * test_na).round())

fill_gender(male_train, female_train, male_test, female_test)

# print(male_train)
# print(male_test)
# print(female_train)
# print(female_test)

print(data_train['Gender'].value_counts())
print(data_test['Gender'].value_counts())
```

```
Male      489
Female    112
Name: Gender, dtype: int64
13
11
Male      500
Female    114
Name: Gender, dtype: int64
Male      295
Female     72
Name: Gender, dtype: int64
```

```
In [47]: se_train_na = data_train.isnull().sum()['Self_Employed']
se_test_na = data_test.isnull().sum()['Self_Employed']

se_sum = data_train['Self_Employed'].value_counts()
se_yes_ratio = se_sum['Yes']/ se_sum.values.sum()
se_no_ratio = se_sum['No']/ se_sum.values.sum()

# print(se_yes_ratio)
# print(se_no_ratio)

def fill_self_employed(train_yes, test_yes, train_no, test_no):
    data_train['Self_Employed'].fillna('Yes', inplace= True, limit = train_yes)
    data_train['Self_Employed'].fillna('No', inplace= True, limit = train_no)
    data_test['Self_Employed'].fillna('Yes', inplace= True, limit = test_yes)
    data_test['Self_Employed'].fillna('No', inplace= True, limit = test_no)

train_yes = int(round(se_yes_ratio * se_train_na))
train_no = int(round(se_no_ratio * se_train_na))
test_yes = int(round(se_yes_ratio * se_test_na))
test_no = int(round(se_no_ratio * se_test_na))
# print(train_yes)
# print(train_no)
# print(test_yes)
# print(test_no)
fill_self_employed(train_yes, test_yes, train_no, test_no)
```

```
In [48]: LAT_train_na = data_train['Loan_Amount_Term'].isnull().sum()
LAT_test_na = data_test['Loan_Amount_Term'].isnull().sum()

LAT_count = data_train['Loan_Amount_Term'].value_counts()
# print(LAT_count)

v_ratio = {}
for i in LAT_count.index:
    v_ratio[i] = LAT_count[i]/ LAT_count.sum()

# (np.array(list(v_ratio.values())) * LAT_train_na).round()
# print(v_ratio)

for ind, v in v_ratio.items():
    v_ratio[ind] = int((v * LAT_train_na).round())

v_ratio_test = {}
for ind, v in v_ratio.items():
    v_ratio_test[ind] = int((v * LAT_test_na).round())

# print(v_ratio_test[480.0])

v360_train = v_ratio[360.0]
v180_train = v_ratio[180.0] + 1
v360_test = v_ratio_test[360.0]
v180_test = v_ratio[180.0]

def fill_Loan_Amount_Term(v360_train, v180_train, v360_test, v180_test):
    data_train['Loan_Amount_Term'].fillna(360.0, inplace= True, limit = v360_train)
    data_train['Loan_Amount_Term'].fillna(180.0, inplace= True, limit = v180_train)
    data_test['Loan_Amount_Term'].fillna(360.0, inplace= True, limit = v360_test)
    data_test['Loan_Amount_Term'].fillna(180.0, inplace= True, limit = v180_test)

fill_Loan_Amount_Term(v360_train, v180_train, v360_test, v180_test)
print(data_train.isnull().sum())
print(data_test.isnull().sum())
```

Loan_ID	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	50
Property_Area	0
Loan_Status	0
dtype: int64	
Loan_ID	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	29
Property_Area	0
dtype: int64	



```
In [49]: CH_na = data_train['Credit_History'].isnull().sum()
CH_test_na = data_test['Credit_History'].isnull().sum()

total = data_train['Credit_History'].value_counts().sum()
one = data_train['Credit_History'].value_counts()[1.0]
zero = data_train['Credit_History'].value_counts()[0]

one_ratio = one/ total
zero_ratio = zero/ total

one_train = int((one_ratio * CH_na).round())
zero_train = int((zero_ratio * CH_na).round())
one_test = int((one_ratio * CH_test_na).round())
zero_test = int((zero_ratio * CH_test_na).round())

def fill_Credit_History(one_train, zero_train, one_test, zero_test):
    data_train['Credit_History'].fillna(1.0, inplace= True, limit = one_train)
    data_train['Credit_History'].fillna(0, inplace= True, limit = zero_train)
    data_test['Credit_History'].fillna(1.0, inplace= True, limit = one_test)
    data_test['Credit_History'].fillna(0, inplace= True, limit = zero_test)

fill_Credit_History(one_train, zero_train, one_test, zero_test)

print(data_train.isnull().sum())
print(data_test.isnull().sum())
```

```
Loan_ID      0
Gender       0
Married      0
Dependents   0
Education    0
Self_Employed  0
ApplicantIncome  0
CoapplicantIncome  0
LoanAmount   0
Loan_Amount_Term  0
Credit_History  0
Property_Area  0
Loan_Status  0
dtype: int64
Loan_ID      0
Gender       0
Married      0
Dependents   0
Education    0
Self_Employed  0
ApplicantIncome  0
CoapplicantIncome  0
LoanAmount   0
Loan_Amount_Term  0
Credit_History  0
Property_Area  0
dtype: int64
```

```
In [50]: data_train['Gender'].replace('Male', 1, inplace = True)
data_train['Gender'].replace('Female', 2, inplace = True)
data_test['Gender'].replace('Male', 1, inplace = True)
data_test['Gender'].replace('Female', 2, inplace = True)

data_train['Married'].replace('Yes', 1, inplace = True)
data_train['Married'].replace('No', 0, inplace = True)
data_test['Married'].replace('Yes', 1, inplace = True)
data_test['Married'].replace('No', 0, inplace = True)

data_train['Self_Employed'].replace('Yes', 1, inplace = True)
data_train['Self_Employed'].replace('No', 0, inplace = True)
data_test['Self_Employed'].replace('Yes', 1, inplace = True)
data_test['Self_Employed'].replace('No', 0, inplace = True)

data_train['Education'].replace('Graduate', 1, inplace = True)
data_train['Education'].replace('Not Graduate', 0, inplace = True)
data_test['Education'].replace('Graduate', 1, inplace = True)
data_test['Education'].replace('Not Graduate', 0, inplace = True)

data_train['Property_Area'].replace('Urban', 1, inplace = True)
data_train['Property_Area'].replace('Rural', 2, inplace = True)
data_train['Property_Area'].replace('Semiurban', 3, inplace = True)
data_test['Property_Area'].replace('Urban', 1, inplace = True)
data_test['Property_Area'].replace('Rural', 2, inplace = True)
data_test['Property_Area'].replace('Semiurban', 3, inplace = True)

data_train['Loan_Status'].replace('Y', 1, inplace = True)
data_train['Loan_Status'].replace('N', 0, inplace = True)

data_train['Dependents'] = data_train['Dependents'].astype(int)
# print(data_train['Dependents'])
```

```
In [51]: y = data_train['Loan_Status']
X = data_train.iloc[:, 1:-1]

y = np.array(y).reshape(-1,1)
X = np.array(X)

print(y.shape)
print(X.shape)
```

```
(614, 1)
(614, 11)
```

```
In [52]: means = X.mean(axis = 0)
stds = X.std(axis = 0)

X_norm = (X - means)/ stds
X_norm = X_norm[:, (5,8,9)]
print(X_norm.shape)
```

```
(614, 3)
```

```
In [53]: percent_train = .8

def partition(X, y, percent_train):
    idx = np.arange(0,y.shape[0])
    random.seed(1412)
    random.shuffle(idx)
    train_size = int(percent_train * len(idx))
    train_idx = idx[:train_size]
    test_idx = idx[train_size:]
    X_train = X[train_idx]
    y_train = y[train_idx]
    X_test = X[test_idx]
    y_test = y[test_idx]
    return idx, X_train, y_train, X_test, y_test
```

```
In [54]: idx, X_train, y_train, X_test, y_test = partition(X_norm, y, percent_train)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

(491, 3)  
(123, 3)  
(491, 1)  
(123, 1)

```
In [55]: data_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null    object
1   Gender                614 non-null    int64
2   Married               614 non-null    int64
3   Dependents            614 non-null    int64
4   Education             614 non-null    int64
5   Self_Employed         614 non-null    int64
6   ApplicantIncome       614 non-null    int64
7   CoapplicantIncome     614 non-null    float64
8   LoanAmount            614 non-null    float64
9   Loan_Amount_Term      614 non-null    float64
10  Credit_History        614 non-null    float64
11  Property_Area         614 non-null    int64
12  Loan_Status           614 non-null    int64
dtypes: float64(4), int64(8), object(1)
memory usage: 62.5+ KB
```


```
In [56]: NM_model = Logistic_NM()

iterations = 1000
theta_initial = np.ones((X_train.shape[1],1))

nm_theta1, nm_cost1 = NM_model.newtonsMethod(X_train, y_train, theta_initial, iterations)
print("theta:",nm_theta1)

print(nm_cost[0])
plt.plot(nm_cost)
plt.xlabel('iterations')
plt.ylabel('cost')
plt.title('Cost Over iterations')
plt.show()
```

Minimum at iteration: 999  
theta: [[0.20574002]  
[0.12499897]  
[1.38603839]]  
1.499471045823371



```
In [57]: print("Accuracy = ",NM_model.getAccuracy(X_test, y_test, nm_theta1))
```

Accuracy = 84.5528455284553

## The report

Write a brief report covering your experiments (both in lab and take home) and submit the Jupyter notebook via JupyterHub at <https://puffer.cs.ait.ac.th> (<https://puffer.cs.ait.ac.th>) before the next lab.

In your solution, be sure to follow instructions!

**if the data is not easily linearly separable, both Newton's method and Linear Regression do not work well. They give us accuracy of around 57.5%. After polar transformation it seems worked better. Newton's method is faster to converge than linear regression model within 300 iterations.**

**For take home exercise, Newton's method converge in 400 iterations and the accuracy is 84.6%.**